

# Applications of Parallel Computing

# Logistics

Lectures: Mondays 15:00-17:00 PM

Last Lecture: 15:00-18:00 PM, we might have a Dec 9<sup>th</sup> lecture

Instructor: Maryam Mehri Dehnavi, [mmehride@cs.toronto.edu](mailto:mmehride@cs.toronto.edu)

Office hours: Mondays 14:00-15:00 PM

See the course syllabus for links to meetings.

# Course Overview

## **Required Prerequisites:**

CSC 367 Parallel Computing

CSC336/CSC436 Numerical Methods

## **Discussions:**

*Piazza* and *Quercus* for all course material along with important announcements.

The course syllabus is in Quercus, please read it carefully!

# Evaluation

**Presentation:** We will provide a list of papers with presentation dates. Each group (composed of two students) will present 1-2 paper and possibly one tool (a debugger, profiler, etc.) depending on the final class enrolment count.

**Project:** An essential component of this course is the course project. You are expected to work on a cutting-edge research problem (in groups of two). The project should have an experimental component and you are expected to deliver a novel contribution that is a scaled-down version of conference paper. The final project report has to be in a 10-page research paper format and should have a [reproducibility appendix](#) for artifact evaluation.

**Reviews and discussions:** For sessions that you are not presenting, you should prepare a short summary of the paper being presented. Also prepare a list of questions that you would ask the presenter during his or her talk. Contributing in the open discussions of research presentations plays an important role in your final grade.

# Evaluation

| <b>Item</b>             | <b>Grade breakdown</b>   |
|-------------------------|--|
| Project                 | 50%  |
| Paper/tool presentation | 20%  |
| Reviews                 | 15%  |
| Participation           | 15% (you should also participate in discussions for papers you did not write a review for) |

# Important Dates

| <b>Item</b>                                  | <b>Due date</b>  |
|--|--|
| One-page project proposal                    | Tuesday, October 12  |
| Four-page mid-term project report            | Friday, November 5   |
| Final project presentation                   | Monday, December 6   |
| Final project report                         | Wednesday, December 8  |
| Paper/tool presentation and review deadlines | Will be released after the last day of enrolment. Check announcements in Quercus and Piazza. |

# Parallel Computing Resources

You can use any of the below for your project

- [SciNet](#): Canada's largest supercomputing center, we will use their teach cluster.
- Amazon Web Services: Have to use your free education credits.
- Any compute resource provided by your research group.

More information to follow in upcoming announcements...

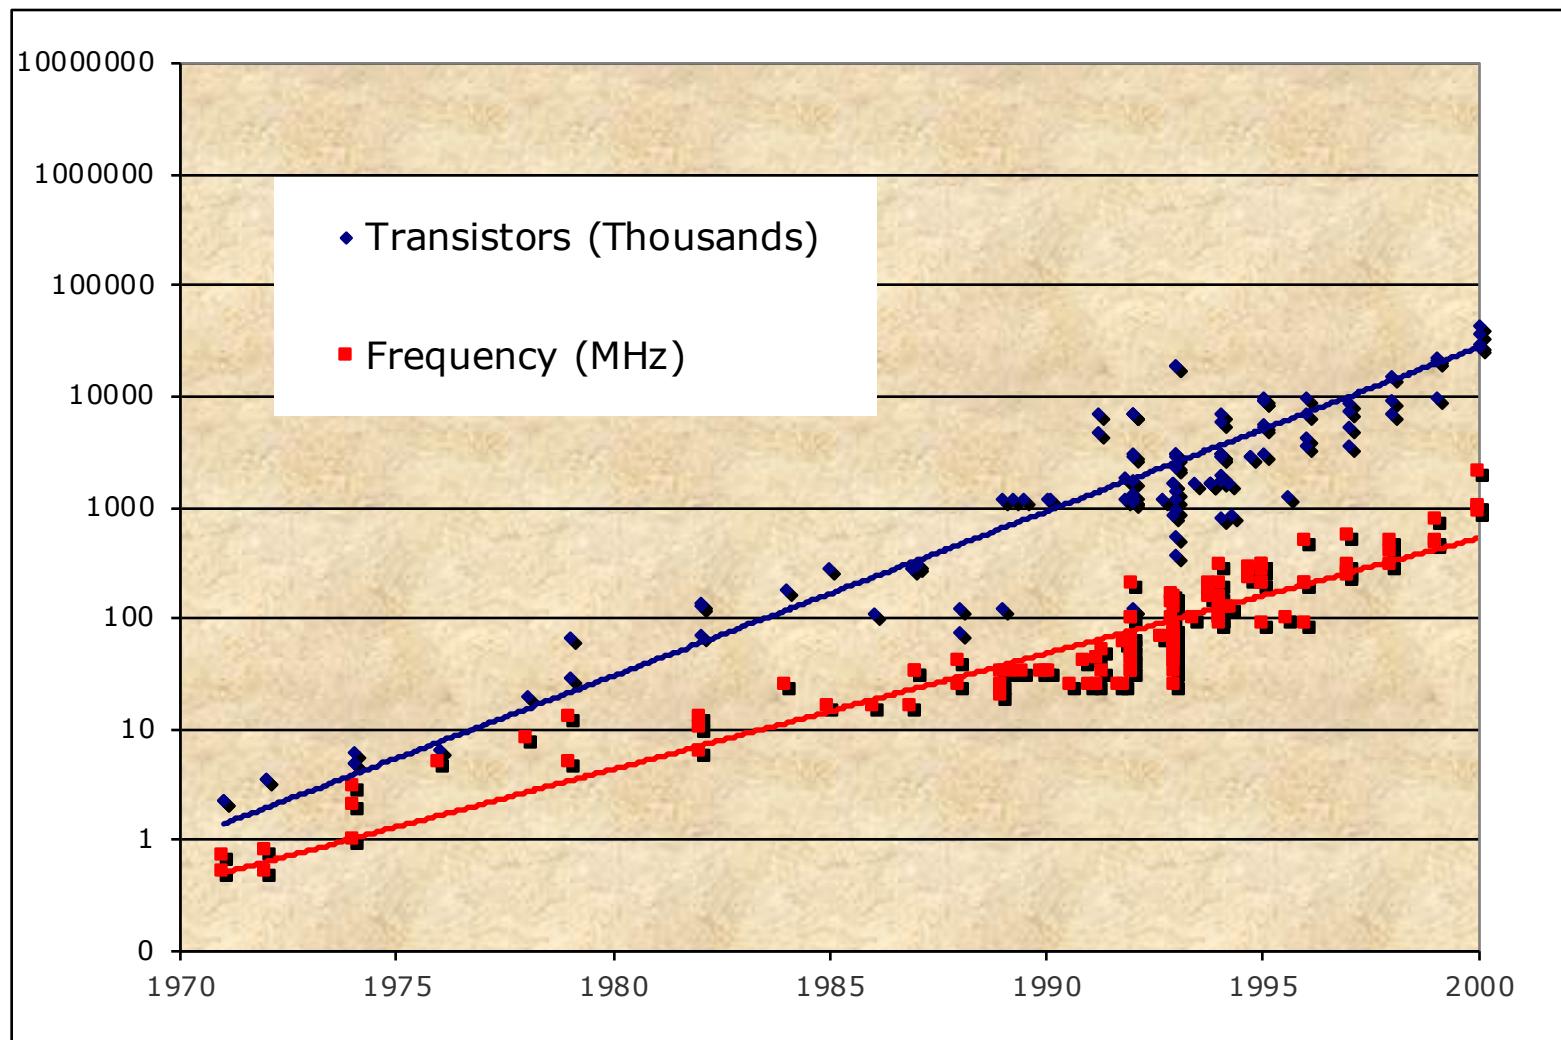
# Tentative Topics

- Overview of parallel architectures and programming models
- The computation Dwarfs in parallel algorithms
- Cloud computing
- Performance models
- Domain-specific code generation
- Parallel algorithms
- Parallel computing applications such as machine learning, biology, etc.

# Overview of Parallel Computing

© Some of the slides in this lecture are adopted from slides from Charles Leiserson and James demmel.

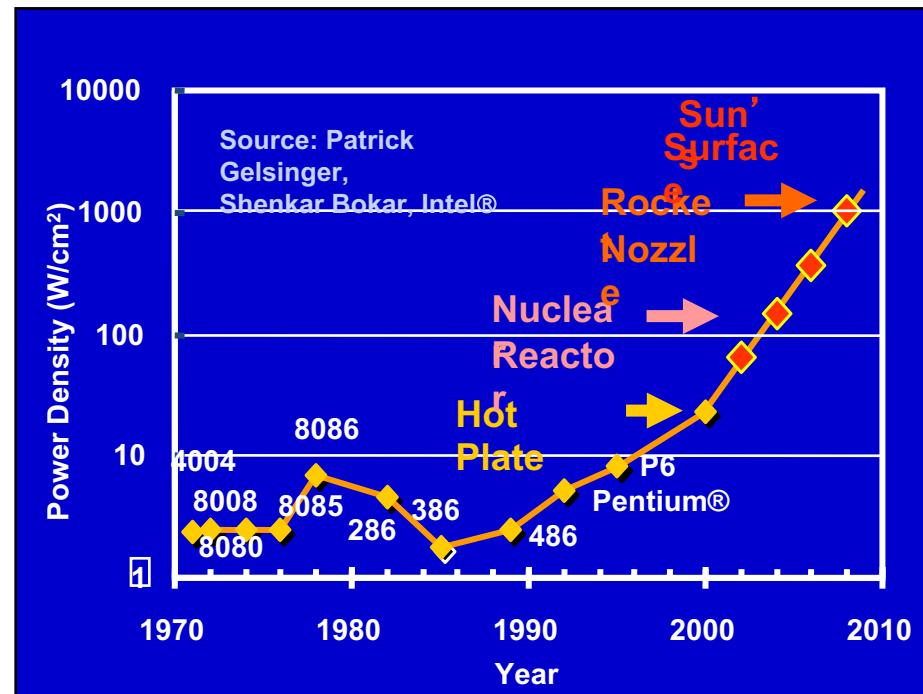
# Microprocessor Transistors / Clock (1970-2000)



2X transistors/Chip Every 1.5 years Called “**Moore’s Law**”

# Power Density Limits Serial Performance

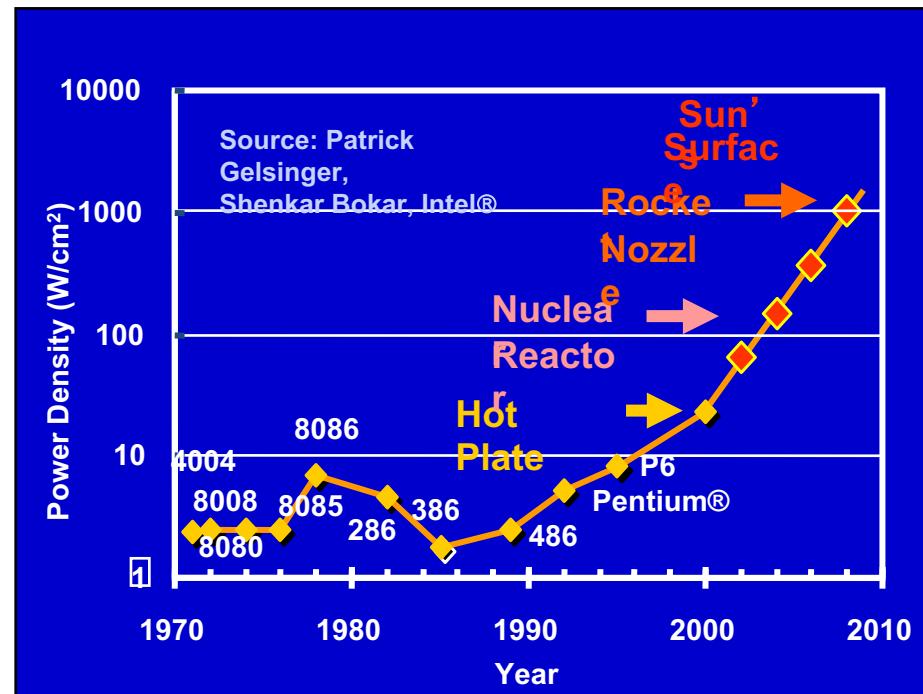
- Concurrent systems are more power efficient
  - Dynamic power is proportional to



Source Demmel

# Power Density Limits Serial Performance

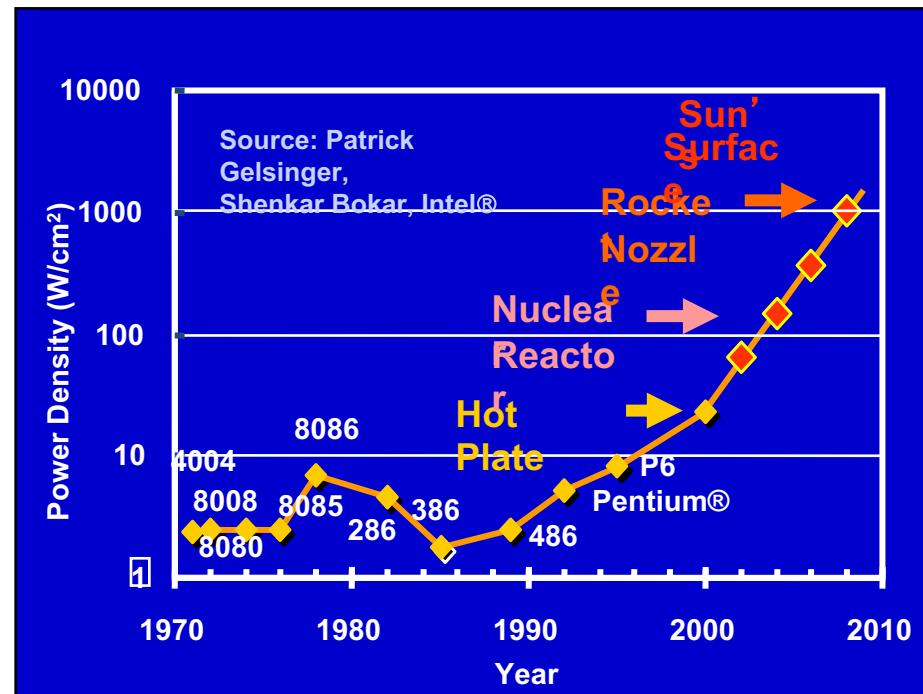
- Concurrent systems are more power efficient
  - Dynamic power is proportional to  $V^2fC$
  - Increasing frequency ( $f$ ) also increases supply voltage ( $V$ ) →



Source Demmel

# Power Density Limits Serial Performance

- Concurrent systems are more power efficient
  - Dynamic power is proportional to  $V^2fC$
  - Increasing frequency ( $f$ ) also increases supply voltage ( $V$ ) → cubic effect
  - Increasing cores increases capacitance ( $C$ ) but only linearly
  - Save power by lowering clock speed
- High performance serial processors waste power
  - Speculation, dynamic dependence checking, etc. burn power
  - Implicit parallelism discovery
- More transistors, but not faster serial processors



Source Demmel

# Revolution in Processors

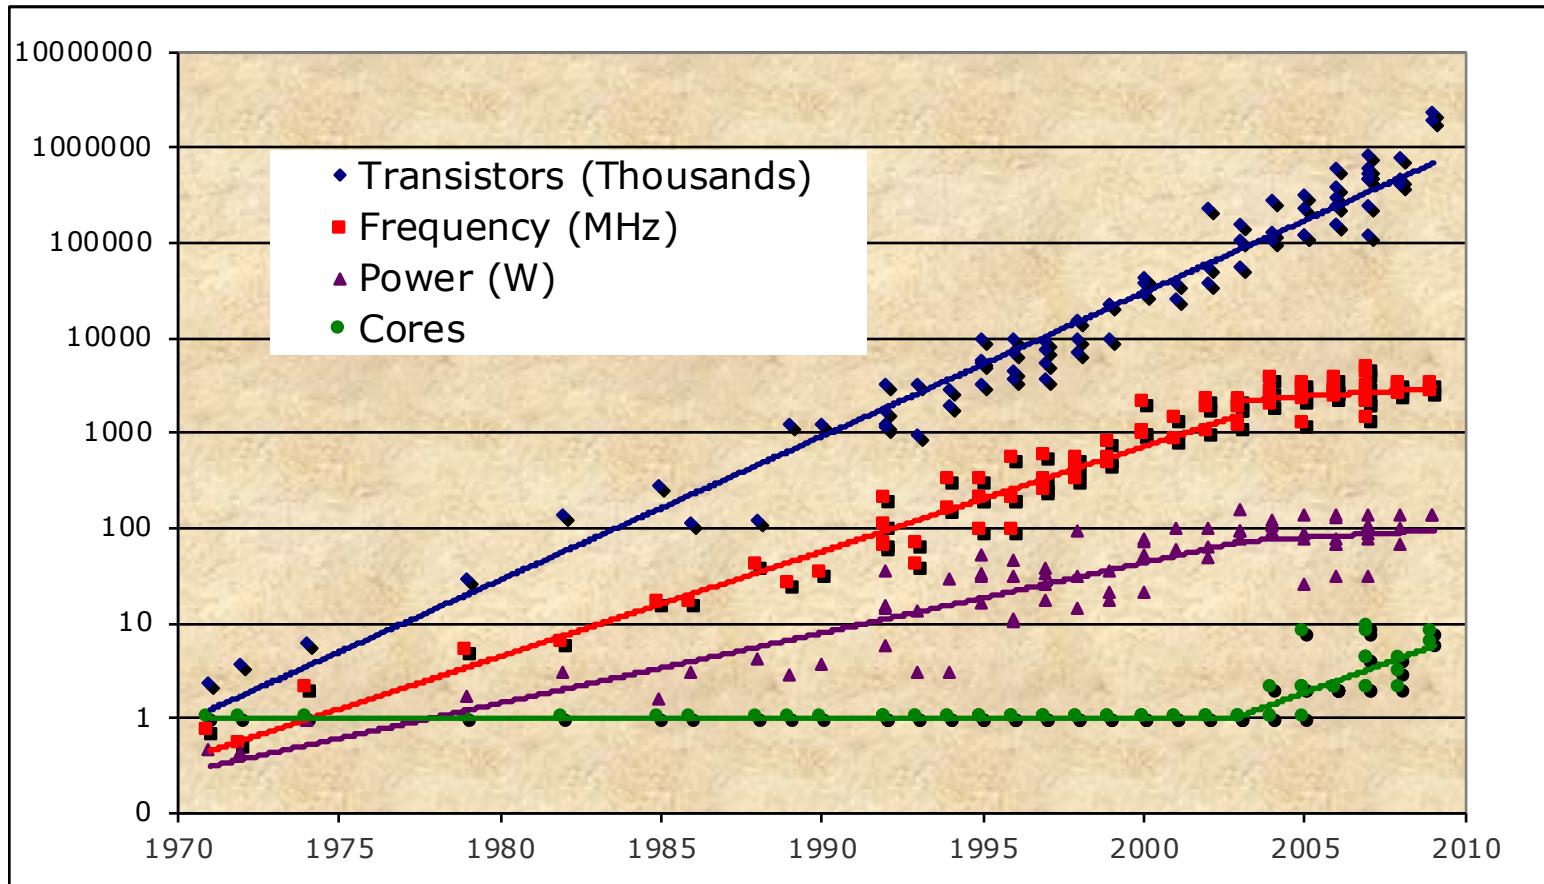
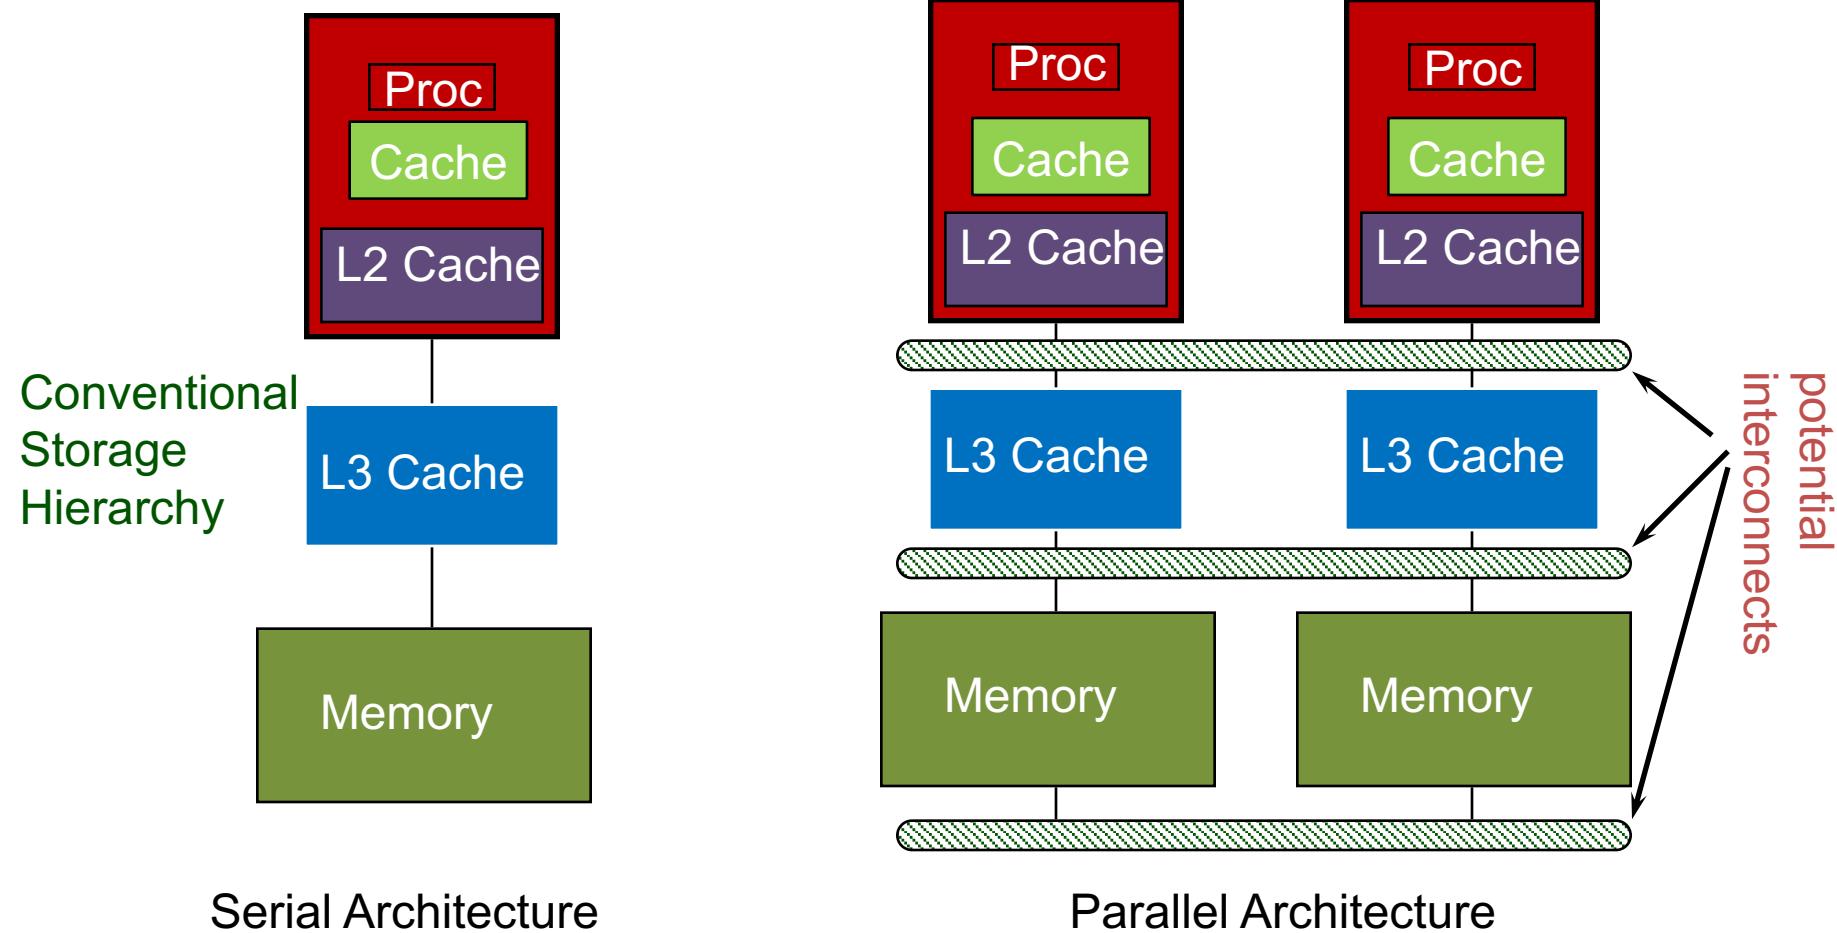


Image source: Leiserson

# Revolution in Processors

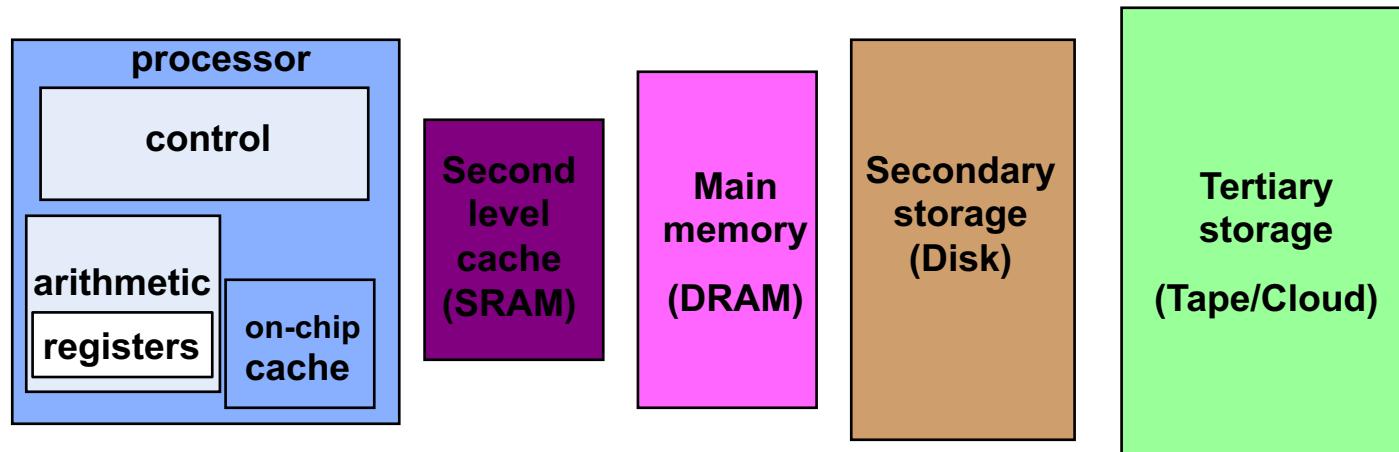
- Chip density is continuing increase ~2x every 2 years
- Clock speed is not
- Number of processor cores may double instead
- Power is under control, no longer growing

# Serial and Parallel Architectures

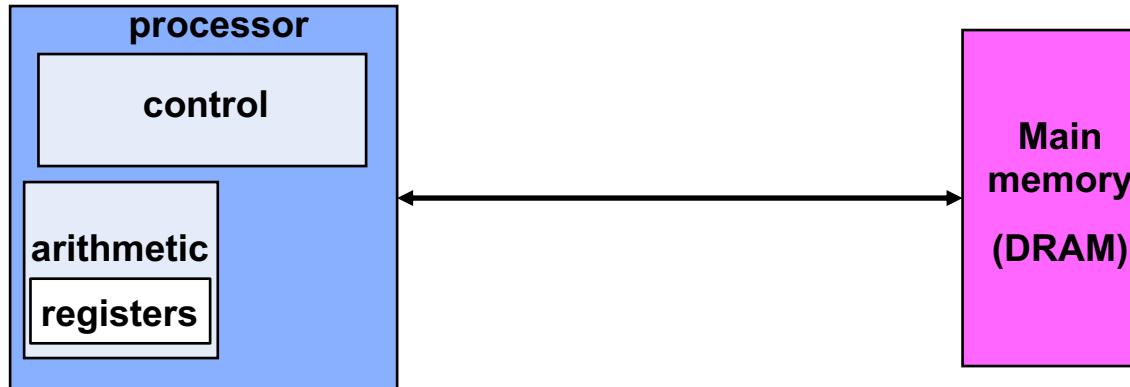


# The Memory Hierarchy

The memory hierarchy of a single processor can contain multiple levels of cache, memory, and storage.



# Bandwidth and Latency



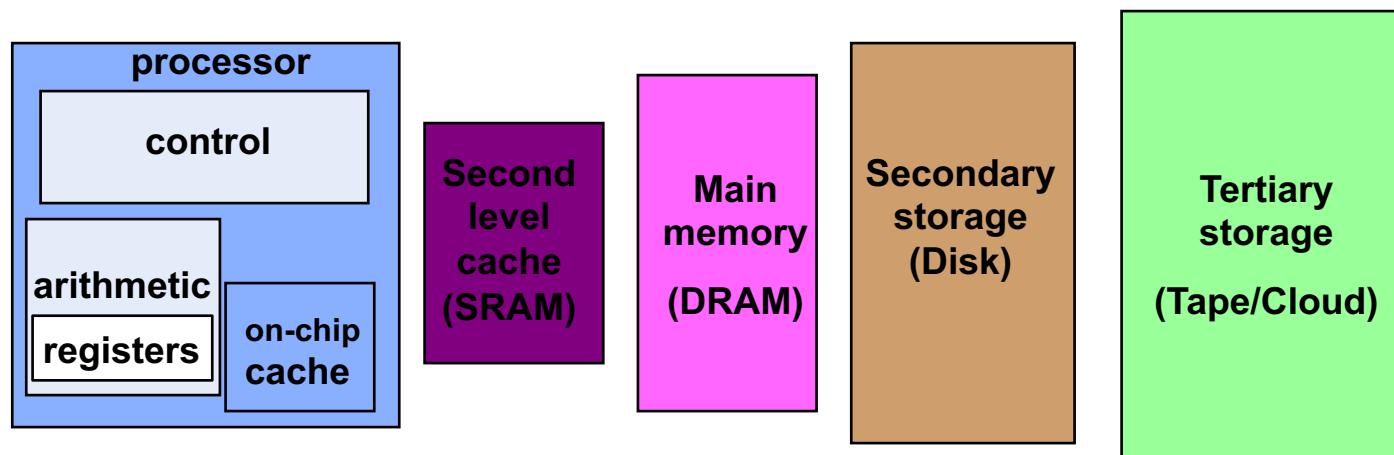
Latency: The time it takes to go from one point to another.

- For example the memory access latency can be 100 ns which means it will take 100ns to retrieve one byte/word from memory.

Bandwidth: The amount of data transferred in a fixed period of time.

- For example the memory bandwidth can be 100MB/s.

# Memory Hierarchy and Latency

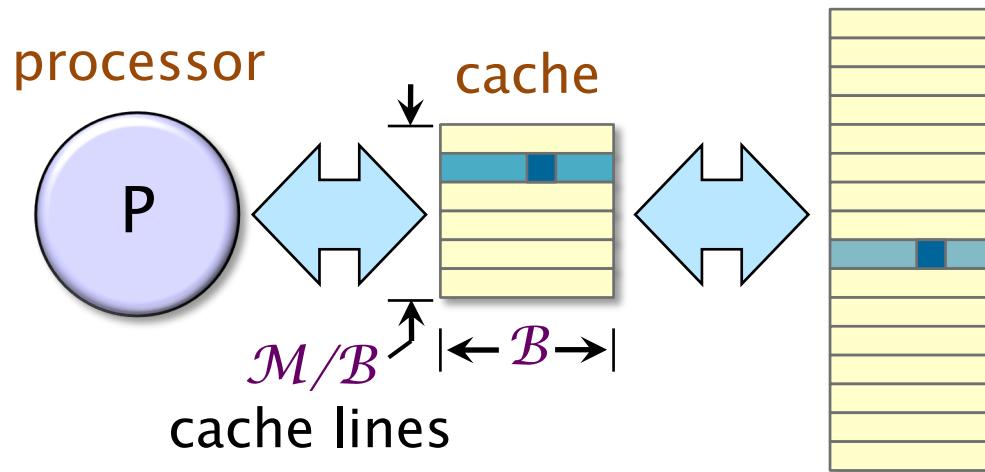


|         |     |      |       |      |       |
|---------|-----|------|-------|------|-------|
| Latency | 1ns | 10ns | 100ns | 10ms | 10sec |
| Size    | KB  | MB   | GB    | TB   | PB    |

# Caches and Locality (Helps hide latency!)

Reads and writes from main memory are in contiguous blocks, i.e. cache lines.

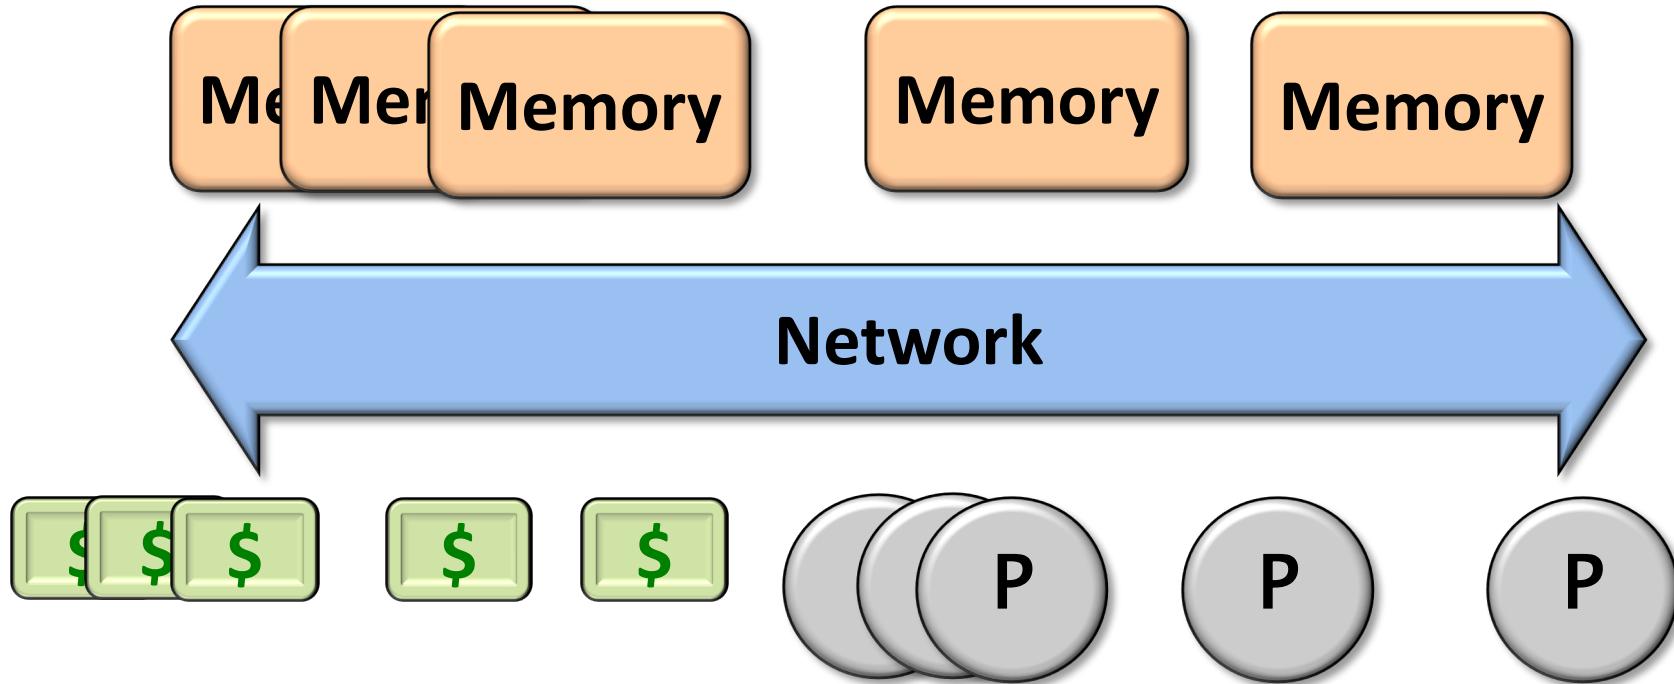
- *Caches* store previously accessed cache lines to hopefully reduce references to memory!
- If the data required by the processor resides in cache we get a *cache hit*, fast!
- Data accesses that are not in the cache lead to a *cache miss*, slow!



# Approaches to Handling Memory Latency

- **Temporal locality:** Reuse data that is already in cache.
  - Eliminate memory operations by saving values in small, fast memory (cache or registers) and reusing them (bandwidth filtering).
  
- **Spatial locality:** Operate on data that are stored close to each other in memory and can be brought into cache together.
  - Take advantage of better bandwidth by getting a chunk of memory into cache (or registers) and using the whole chunk.

# Essential Components of Parallel Architectures



Where is the memory physically located?  
Is it connected directly to processors?  
What is the connectivity of the network?

# Parallel Machine Models

?

# Parallel Machine Models

Shared Memory

Distributed Memory

SIMD and Vector

Hybrid

# Parallel Machine Models

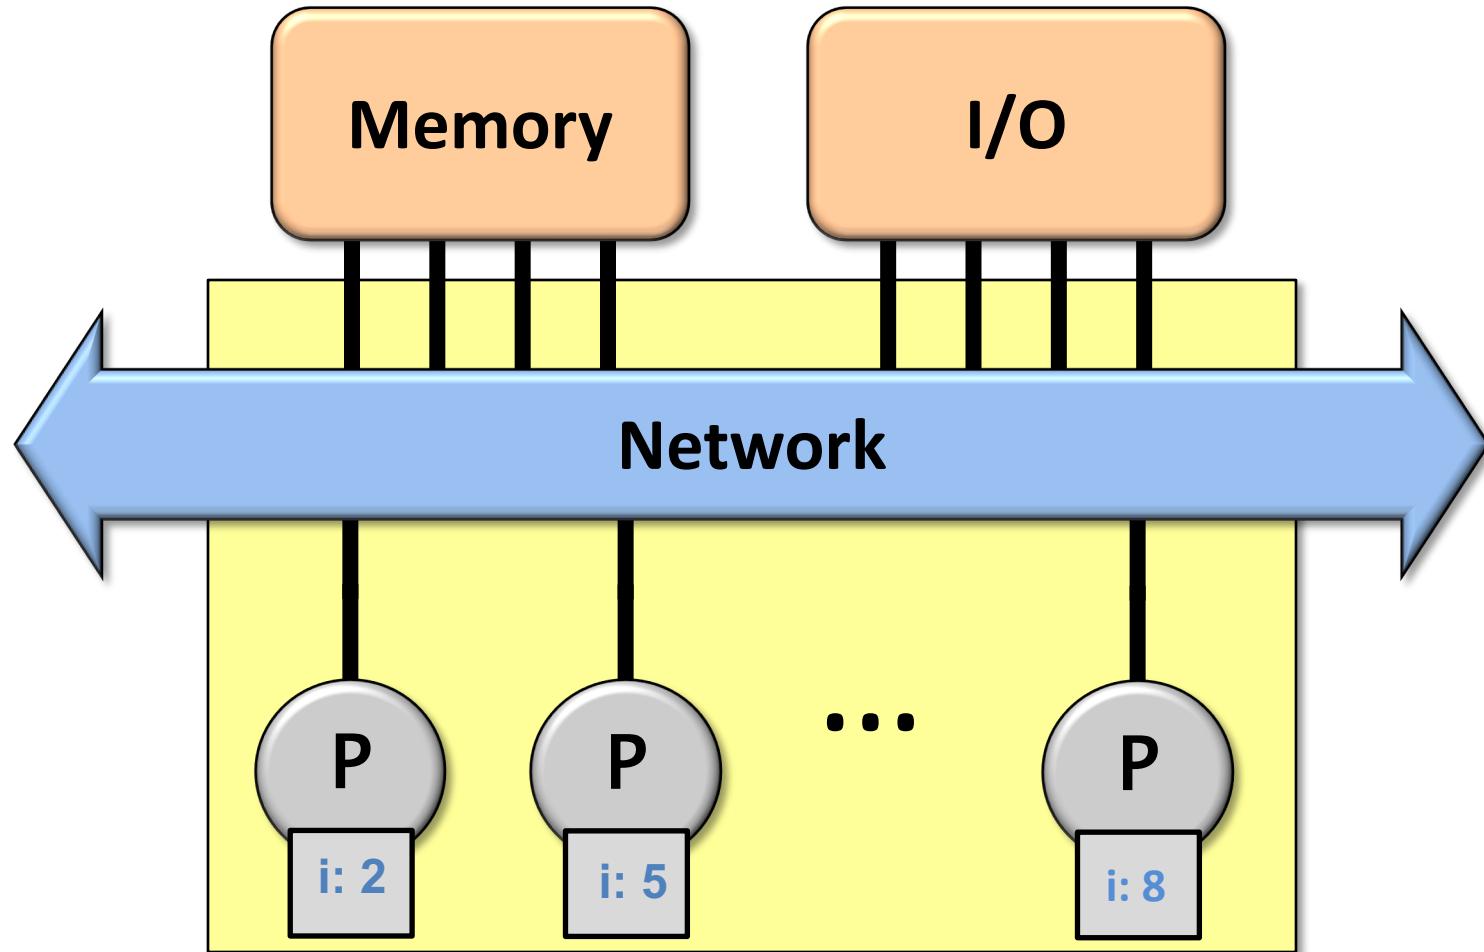
**Shared Memory:** chip multiprocessors,  
multicore architectures, etc.

Distributed Memory

SIMD and Vector: GPUs, DSPs, etc.

Hybrid

# Shared Memory Architecture



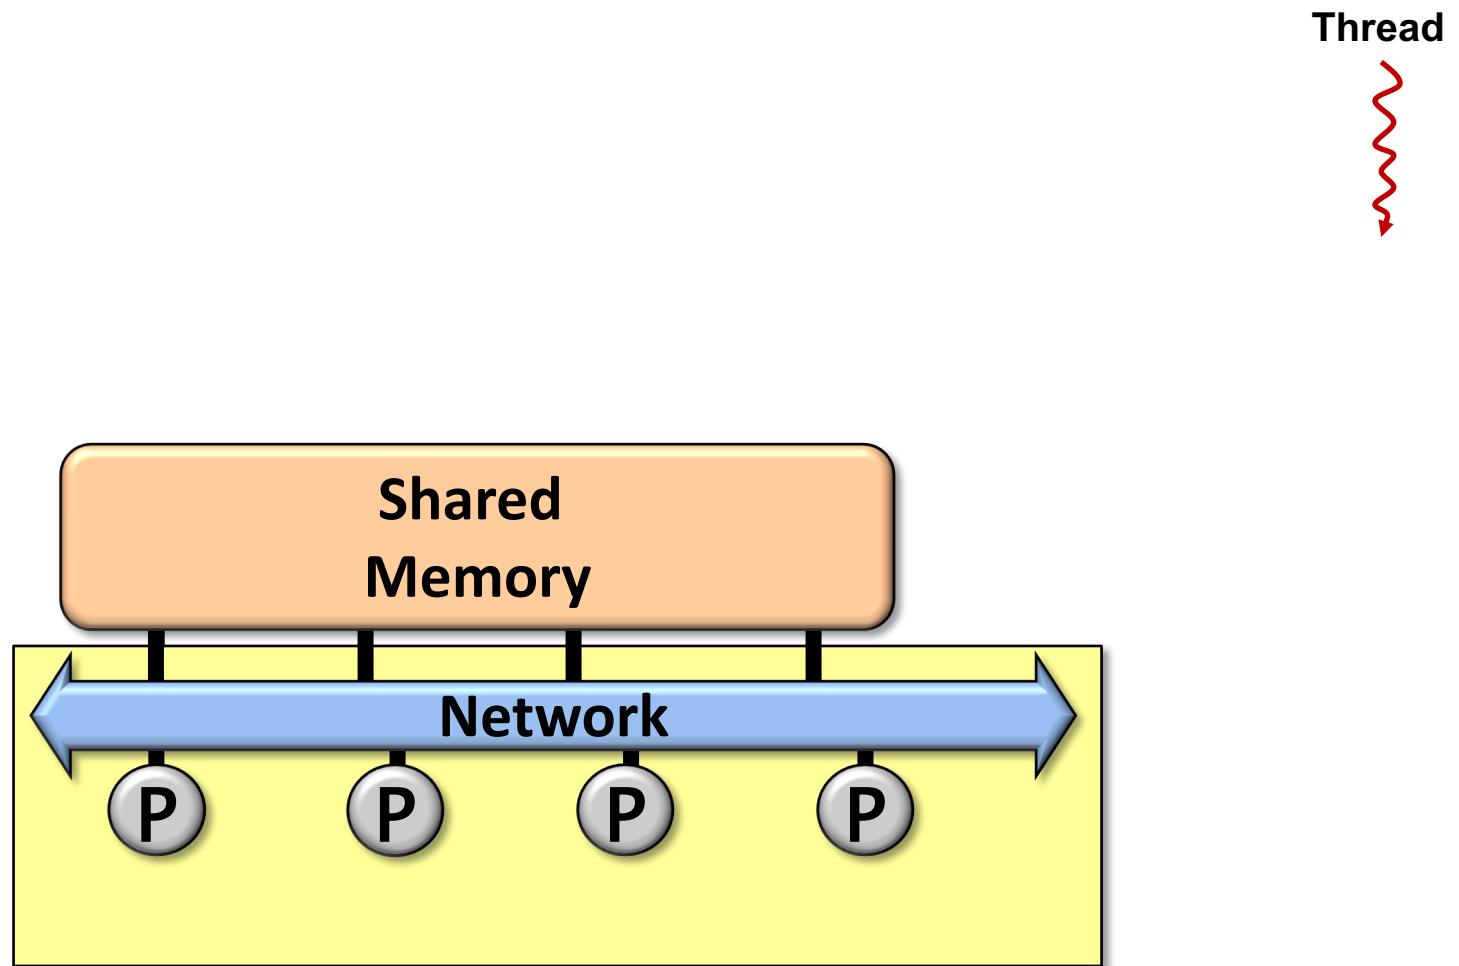
**Chip Multiprocessor (CMP)**

# Parallel Programming Models

- Programming model is made up of the languages and libraries that create an abstract view of the machine
- Control
  - How is parallelism created?
  - What orderings exist between operations?
- Data
  - What data is private vs. shared?
  - How is logically shared data accessed or communicated?
- Synchronization
  - What operations can be used to coordinate parallelism?
  - What are the atomic (indivisible) operations?

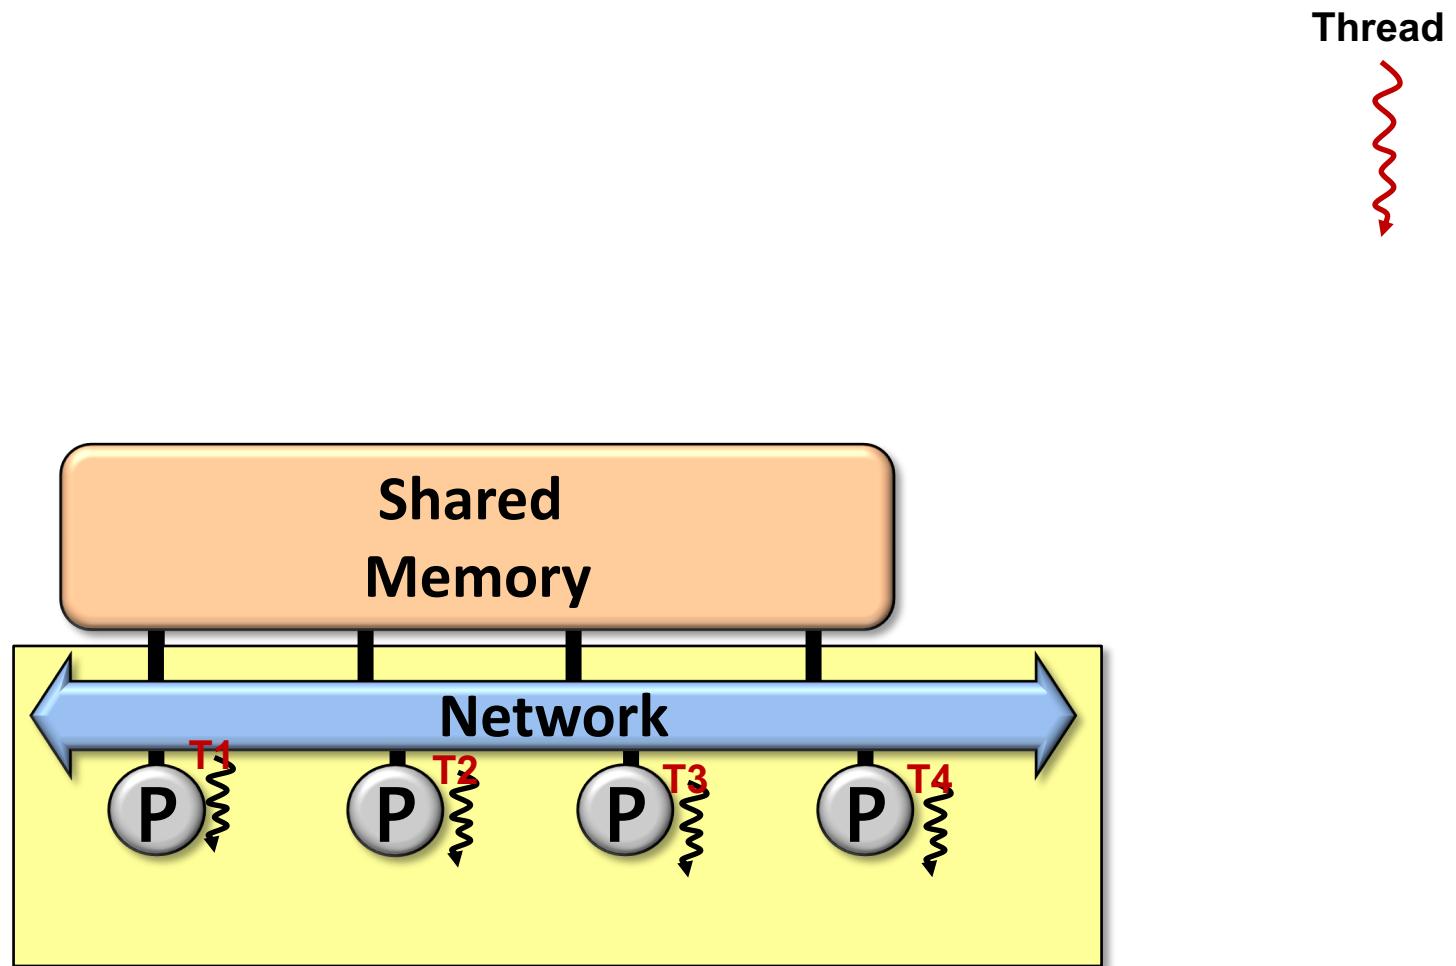
# Programming Model: Shared Memory

Program is a collection of threads of control, can be created mid-execution.



# Programming Model: Shared Memory

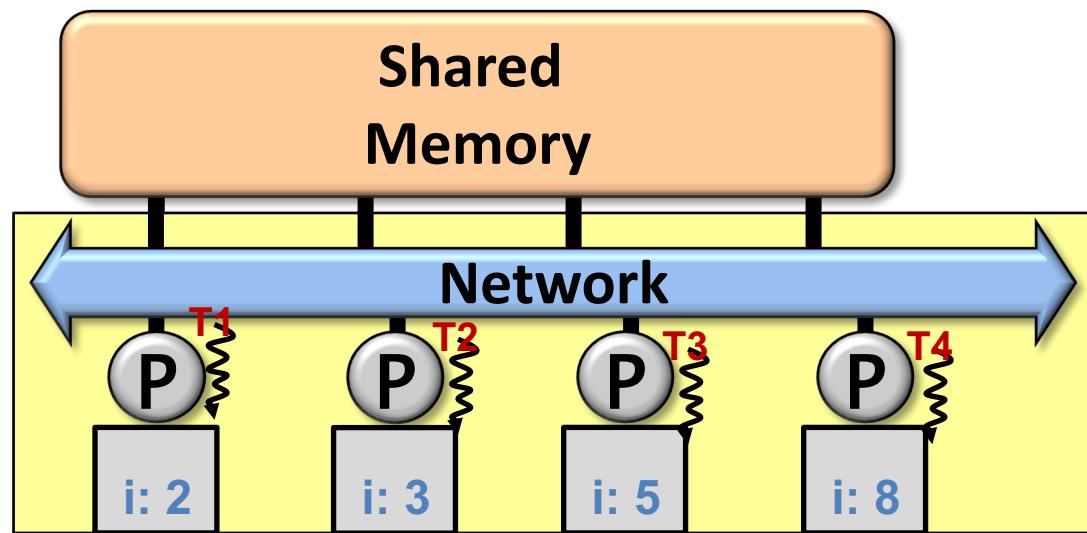
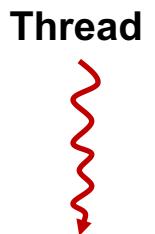
Program is a collection of threads of control, can be created mid-execution.



# Programming Model: Shared Memory

Program is a collection of threads of control, can be created mid-execution.

Each thread has a set of **private variables**, e.g., local stack variables.

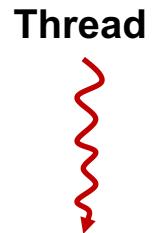


# Programming Model: Shared Memory

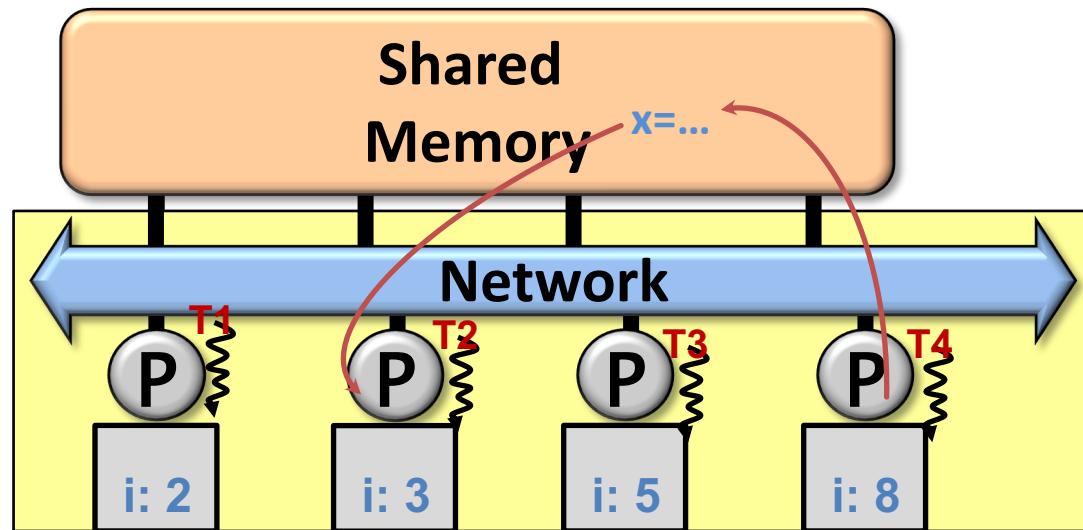
Program is a collection of threads of control, can be created mid-execution.

Each thread has a set of **private variables**, e.g., local stack variables.

Also a set of **shared variables**, e.g., static variables.



- Threads communicate implicitly by writing and reading shared variables.



# Shared memory: Challenges

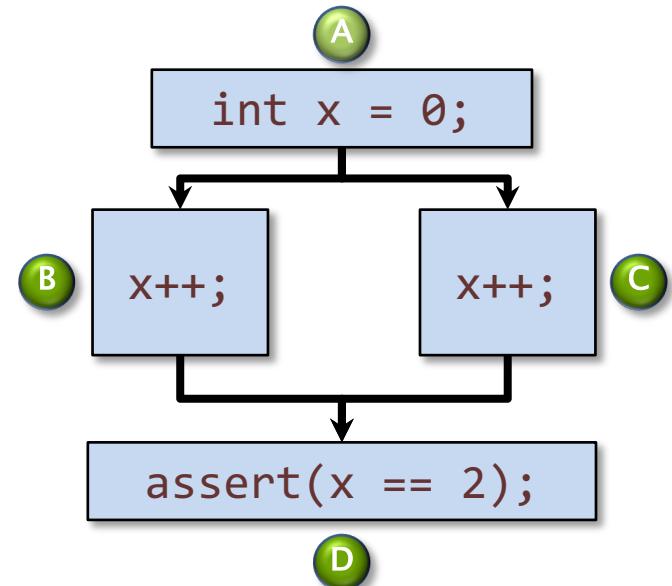
- Races and Locks

# Races

A **race** occurs when two logically parallel instructions access the same memory location and at least one of the instructions performs a write.

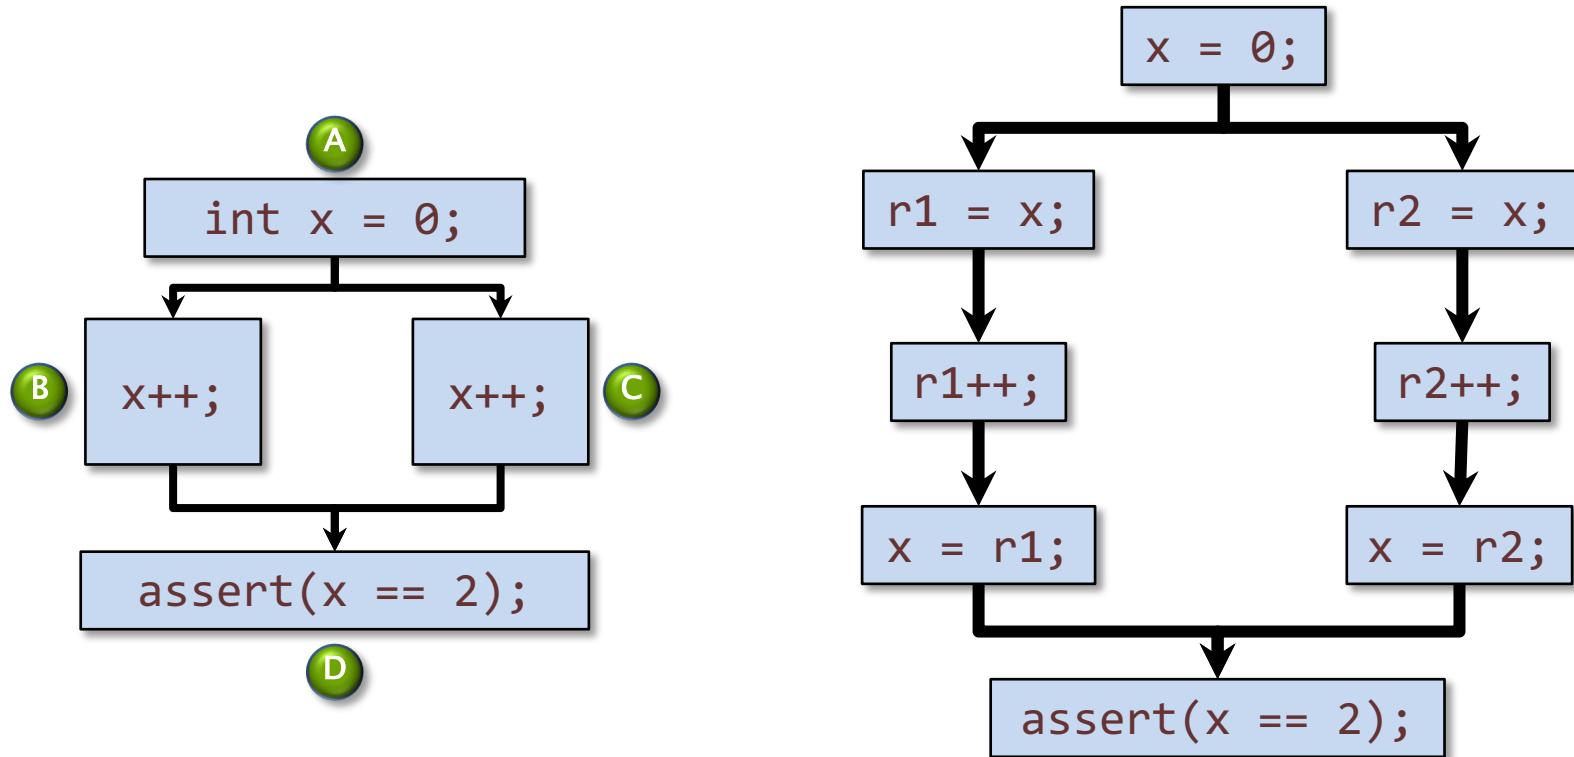
## Example

```
A int x = 0;  
B parallel_for (int i=0, i<2, ++i) {  
C     x++;  
D }  
E assert(x == 2);
```



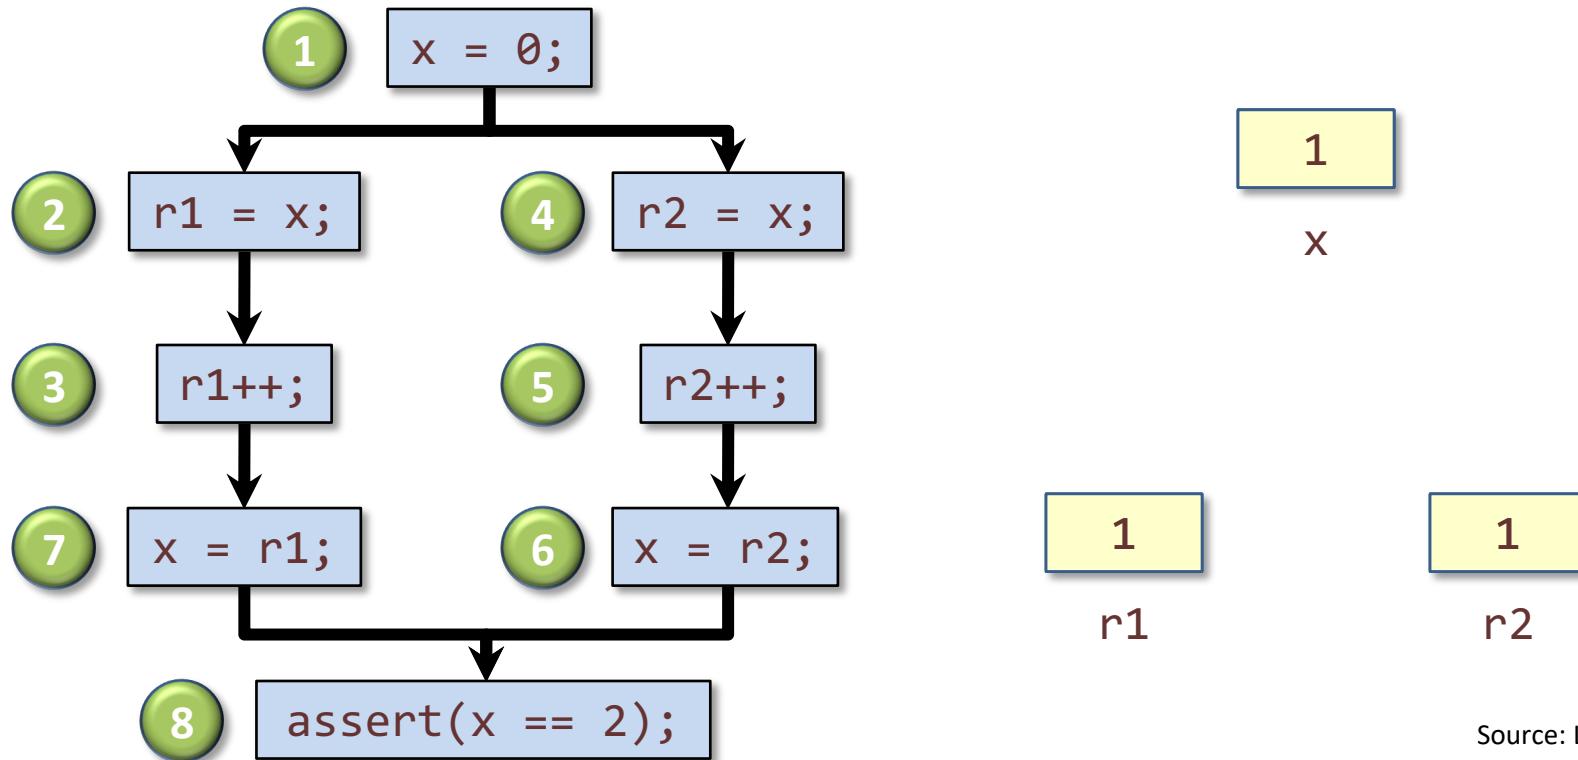
**dependency graph**

# A Closer Look



# Race Bugs

A **race** occurs when two logically parallel instructions access the same memory location and at least one of the instructions performs a write.



Source: Leiserson

# Simple Example

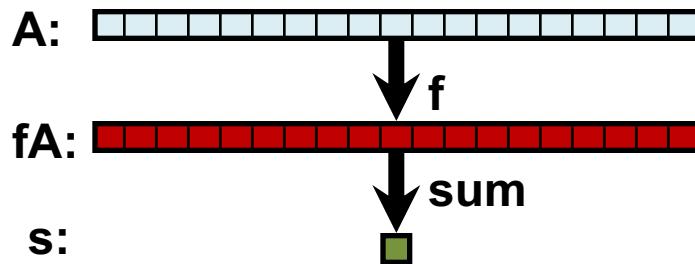
- Consider applying a function  $f$  to the elements of an array  $A$  and then computing its sum:

$$\text{FIR filter: } y[n] = \sum_{i=0}^{n-1} b_i x[n - i]$$

$$\sum_{i=0}^{n-1} f(A[i])$$

- Questions:
  - Where does  $A$  live? All in single memory? Partitioned?

**$A$  = array of all data**  
 **$fA = f(A)$**   
 **$s = \text{sum}(fA)$**



# Simple Example

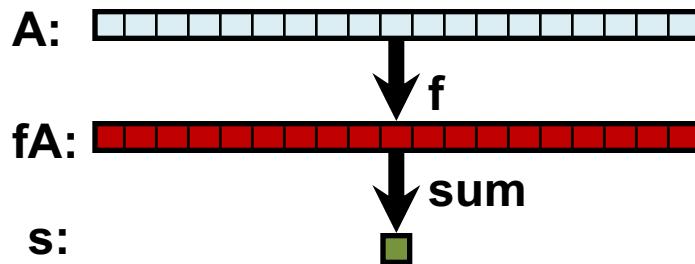
- Consider applying a function  $f$  to the elements of an array  $A$  and then computing its sum:

$$\text{FIR filter: } y[n] = \sum_{i=0}^{n-1} b_i x[n - i]$$

$$\sum_{i=0}^{n-1} f(A[i])$$

- Questions:
  - Where does  $A$  live? All in single memory? Partitioned?

**$A$  = array of all data**  
 **$fA = f(A)$**   
 **$s = \text{sum}(fA)$**



# Shared Memory Code for Computing a Sum

A= 

|   |   |
|---|---|
| 3 | 5 |
|---|---|

 f(x) = x<sup>2</sup>

static int s = 0;

Thread 1

```
for i = 0, n/2-1  
s = s + f(A[i])
```

Thread 2

```
for i = n/2, n-1  
s = s + f(A[i])
```

- What is the problem with this program?

# Shared Memory Code for Computing a Sum

A= 

|   |   |
|---|---|
| 3 | 5 |
|---|---|

 f(x) = x<sup>2</sup>

static int s = 0;

## Thread 1

....

compute  $f([A[i]])$  and put in reg0

9

reg1 = s

0

reg1 = reg1 + reg0

9

s = reg1

9

...

## Thread 2

...

compute  $f([A[i]])$  and put in reg0

25

reg1 = s

0

reg1 = reg1 + reg0

25

s = reg1

25

...

- Assume  $A = [3, 5]$ ,  $f(x) = x^2$ , and  $s=0$  initially
- For this program to work,  $s$  should be  $3^2 + 5^2 = 34$  at the end
  - but it may be 34, 9, or 25

# Improved Code for Computing a Sum

```
static int s = 0;  
static lock lk;
```

**Why not do lock  
Inside loop?**

Thread 1

```
local_s1= 0  
for i = 0, n/2-1  
    local_s1 = local_s1 + f(A[i])  
lock(lk);  
    s = s + local_s1  
unlock(lk);
```

Thread 2

```
local_s2 = 0  
for i = n/2, n-1  
    local_s2= local_s2 + f(A[i])  
lock(lk);  
    s = s +local_s2  
unlock(lk);
```

- Since addition is associative, it's OK to rearrange order
- Most computation is on private variables
  - Sharing frequency is also reduced, which might improve speed
  - But there is still a race condition on the update of shared s
  - The race condition can be fixed by adding locks (only one thread can hold a lock at a time; others wait for it)

# Shared Memory Programming

Several Thread Libraries/systems

- PTHREADS is the POSIX Standard
  - Relatively low level
  - Portable but possibly slow; relatively heavyweight
- OpenMP standard for application level programming
  - Support for scientific programming on shared memory
  - [openmp.org](http://openmp.org)
- TBB: Thread Building Blocks
  - Intel
- CILK: Language of the C “ilk”
  - Lightweight threads embedded into C
- Java threads
  - Built on top of POSIX threads
  - Object within Java language

# Example with OpenMP

## Matrix Vector Multiply OpenMP

```
#pragma omp parallel default (none) shared (A, W, Y, m) private  
(i,j,sum) num_threads(4)  
  
for(i=0; i < m; i++)  
  
{ sum = 0.0;  
  for(j=0; j < m; j++)  
    sum += A[i][j]*W[j];  
  Y[i] =sum;  
}
```

# Parallel Machine Models

Shared Memory: chip multiprocessors,  
multicore architectures, etc.

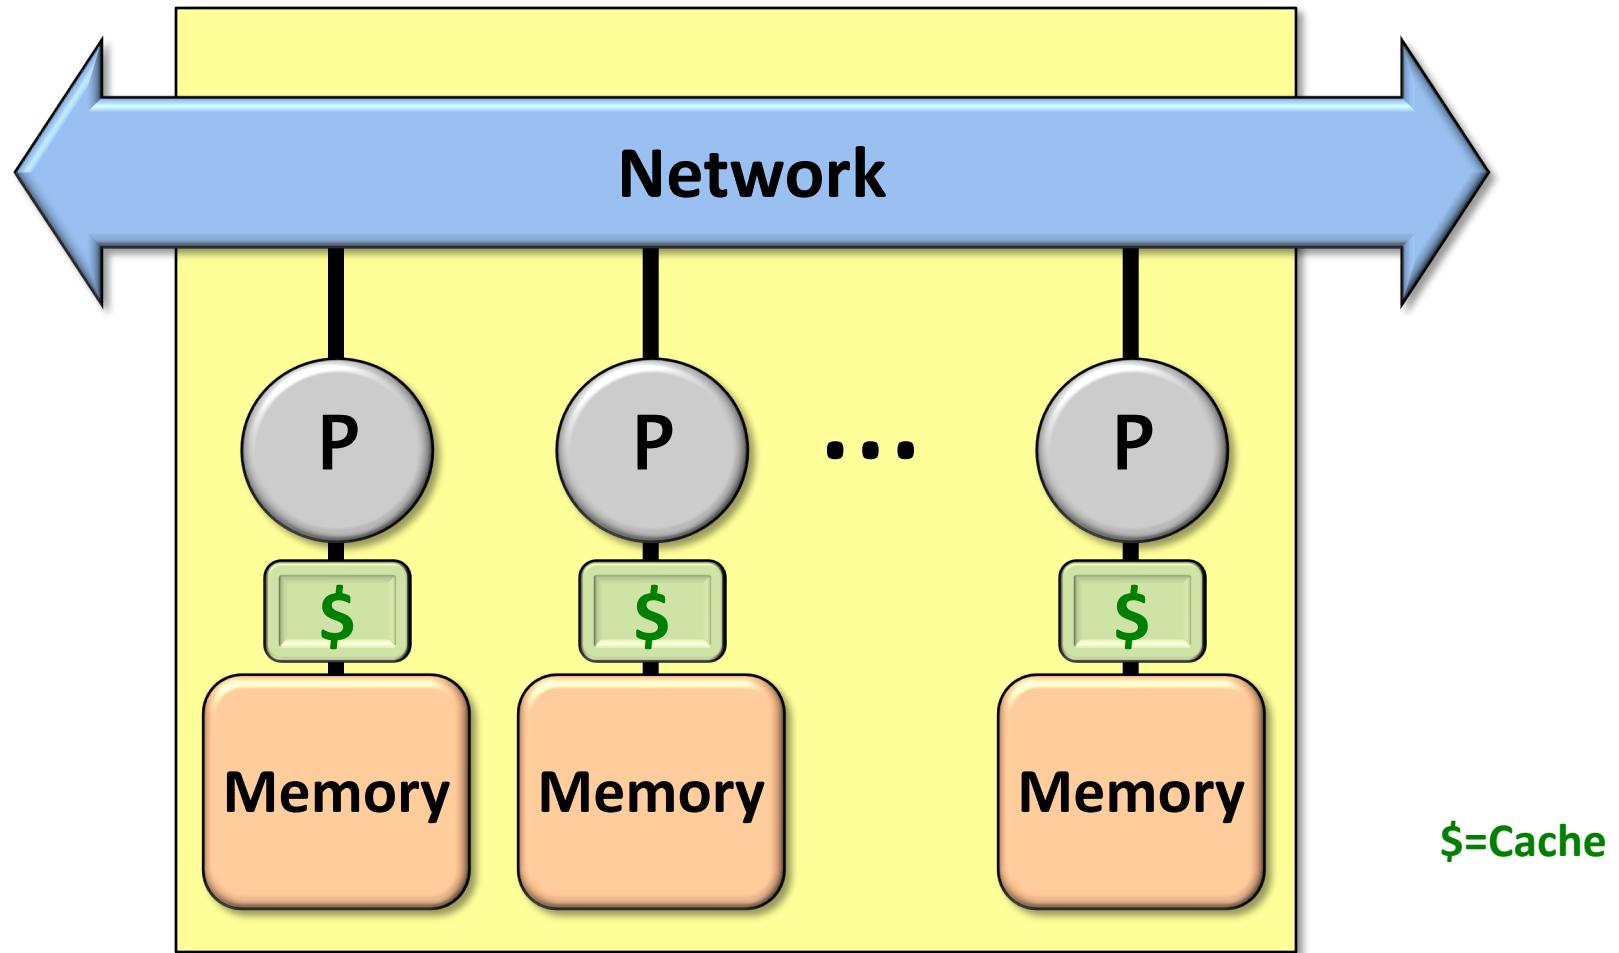
Distributed Memory

SIMD and Vector: GPUs, DSPs, etc.

Hybrid

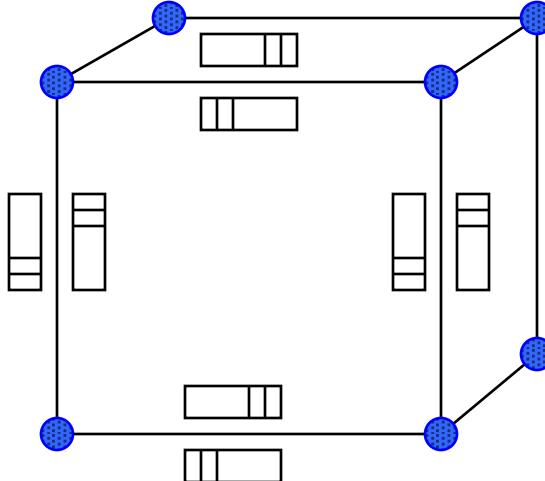
# Distributed Memory Machines and Programming

# Distributed Memory Architecture



# Historical Perspective

- Early distributed memory machines were:
  - Collection of microprocessors.
  - Communication was performed using bi-directional queues between nearest neighbors.
- Messages were forwarded by processors on path.
  - “Store and forward” networking
- There was a strong emphasis on topology in algorithms, in order to minimize the number of hops = minimize time

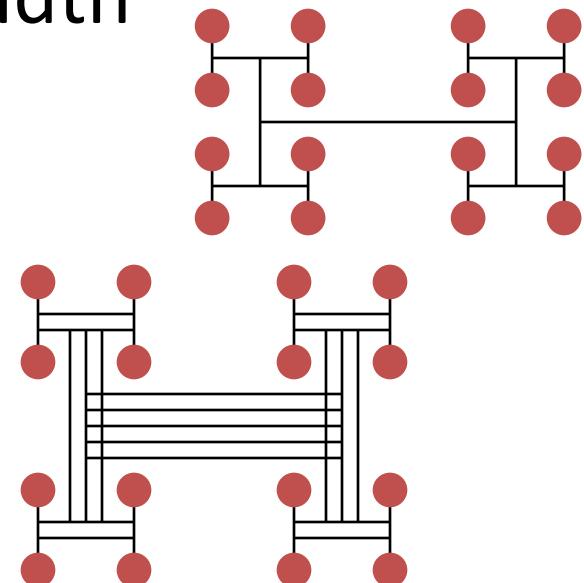
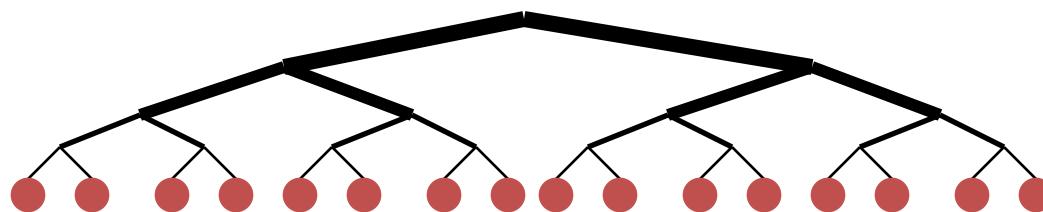


# Design Characteristics of a Network

- **Topology** (how things are connected)
  - Crossbar; ring; 2-D, 3-D, higher-D mesh or torus; hypercube; tree; butterfly; perfect shuffle, dragon fly, ...
- **Routing algorithm:**
  - Example in 2D torus: all east-west then all north-south (avoids deadlock).
- **Switching strategy:**
  - Circuit switching: full path reserved for entire message, like the telephone.
  - Packet switching: message broken into separately-routed packets, like the post office, or internet
- **Flow control** (what if there is congestion):
  - Stall, store data temporarily in buffers, re-route data to other nodes, tell source node to temporarily halt, discard, etc.

# Example Network Topology: Trees

- Many tree algorithms
- Fat trees avoid bisection bandwidth problem:
  - More (or wider) links near top.



# Programming Distributed Memory Machines with Message Passing

# Message Passing Libraries

- Many “message passing libraries” were once available
  - Chameleon, from ANL.
  - CMMD, from Thinking Machines.
  - Express, commercial.
  - MPL, native library on IBM SP-2.
  - NX, native library on Intel Paragon.
  - Zipcode, from LLL.
  - PVM, Parallel Virtual Machine, public, from ORNL/UTK.
  - Others...
  - MPI, Message Passing Interface, now the industry standard.
- Need standards to write portable code.

# Finding Out About the Environment

- Two important questions that arise early in a parallel program are:
  - How many processes are participating in this computation?
  - Which one am I?
- MPI provides functions to answer these questions:
  - **`MPI_Comm_size`** reports the number of processes.
  - **`MPI_Comm_rank`** reports the *rank*, a number between 0 and size-1, identifying the calling process

# Hello (C)

```
#include "mpi.h"
#include <stdio.h>

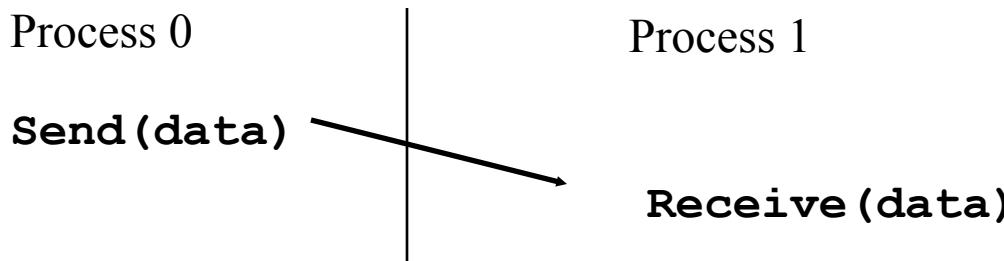
int main( int argc, char *argv[] )
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

# Notes on Hello World

- All MPI programs begin with `MPI_Init` and end with `MPI_Finalize`
- `MPI_COMM_WORLD` is defined by `mpi.h` (in C) or `mpif.h` (in Fortran) and designates all processes in the MPI “job”
- Each statement executes independently in each process
  - including the `printf/print` statements
- The MPI-1 Standard does not specify how to run an MPI program, but many implementations provide  
`mpirun –np 4 a.out`

# MPI Basic Send/Receive

- We need to fill in the details in



- Things that need specifying:
  - How will “data” be described?
  - How will processes be identified?
  - How will the receiver recognize/screen messages?
  - What will it mean for these operations to complete?

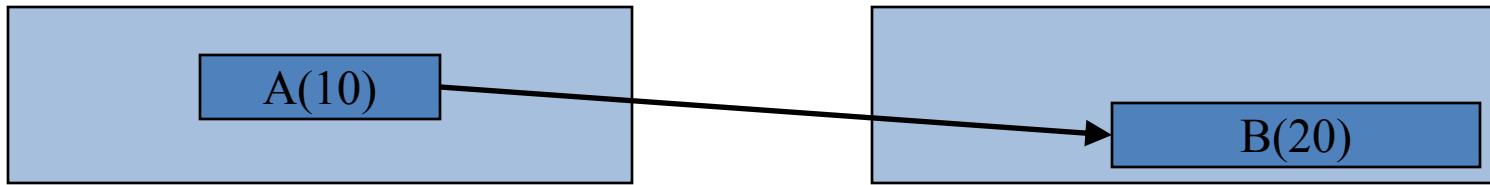
# MPI Datatypes

- The data in a message to send or receive is described by a triple (address, count, datatype), where
- An MPI datatype is defined as:
  - predefined, corresponding to a data type from the language (e.g., MPI\_INT, MPI\_DOUBLE)
  - a contiguous array of MPI datatypes
  - a strided block of datatypes
  - an indexed array of blocks of datatypes
  - an arbitrary structure of datatypes

# MPI Tags

- Messages are sent with an accompanying user-defined integer tag, to assist the receiving process in identifying the message
- Messages can be screened at the receiving end by specifying a specific tag, or not screened by specifying MPI\_ANY\_TAG as the tag in a receive

# MPI Basic (Blocking) Send



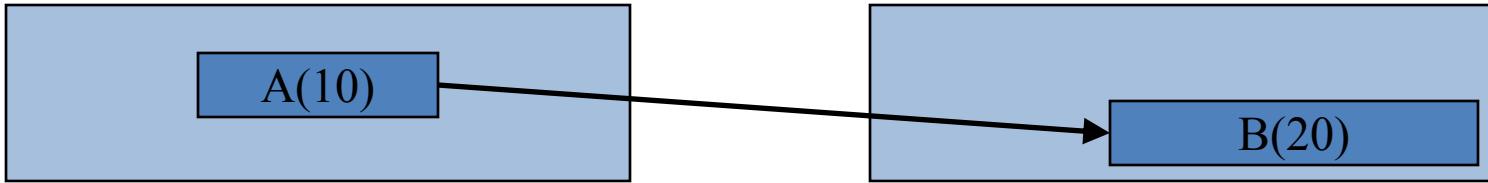
`MPI_Send( A, 10, MPI_DOUBLE, 1, ... )`

`MPI_Recv( B, 20, MPI_DOUBLE, 0, ... )`

## **MPI\_SEND(start, count, datatype, dest, tag, comm)**

- The message buffer is described by **(start, count, datatype)**.
- The target process is specified by **dest**, which is the rank of the target process in the communicator specified by **comm**.
- When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.

# MPI Basic (Blocking) Receive



`MPI_Send( A, 10, MPI_DOUBLE, 1, ... )`

`MPI_Recv( B, 20, MPI_DOUBLE, 0, ... )`

## **`MPI_RECV(start, count, datatype, source, tag, comm, status)`**

- Waits until a matching (both **source** and **tag**) message is received from the system, and the buffer can be used
- **source** is rank in communicator specified by **comm**, or **MPI\_ANY\_SOURCE**
- **tag** is a tag to be matched or **MPI\_ANY\_TAG**
- receiving fewer than **count** occurrences of **datatype** is OK, but receiving more is an error
- **status** contains further information (e.g. size of message)

# A Simple MPI Program

```
#include "mpi.h"
#include <stdio.h>
int main( int argc, char *argv[])
{
    int rank, buf;
    MPI_Status status;
    MPI_Init(&argv, &argc);
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    /* Process 0 sends and Process 1 receives */
    if (rank == 0) {
        buf = 123456;
        MPI_Send( &buf, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
    }
    else if (rank == 1) {
        MPI_Recv( &buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
                  &status );
        printf( "Received %d\n", buf );
    }

    MPI_Finalize();
    return 0;
}
```

# MPI is Simple

- Many parallel programs can be written using just these six functions, only two of which are non-trivial:
  - **MPI\_INIT**
  - **MPI\_FINALIZE**
  - **MPI\_COMM\_SIZE**
  - **MPI\_COMM\_RANK**
  - **MPI\_SEND**
  - **MPI\_RECV**

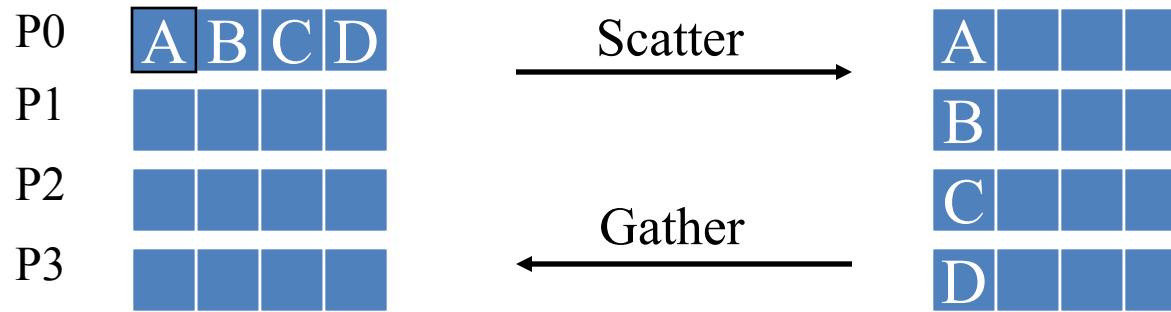
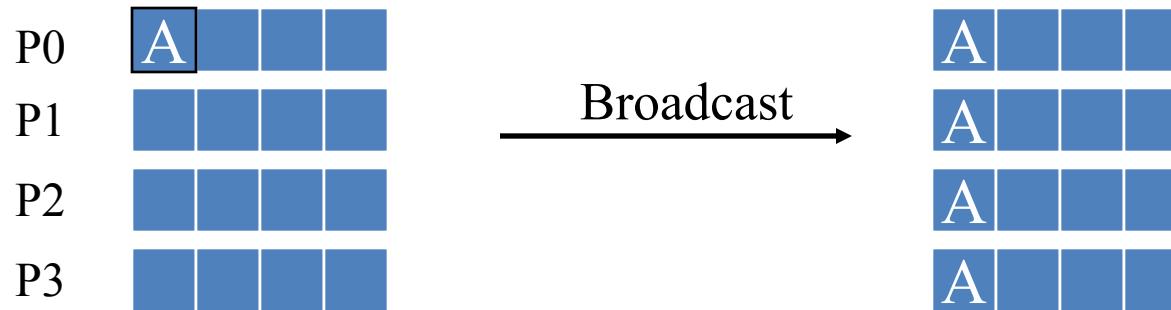
# Another Approach to Parallelism

- *Collective* routines provide a higher-level way to organize a parallel program
- Each process executes the same communication operations
- MPI provides a rich set of collective operations...

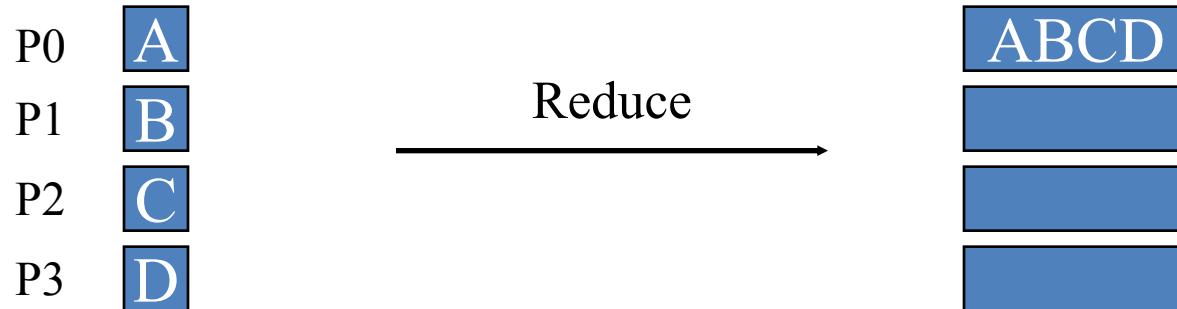
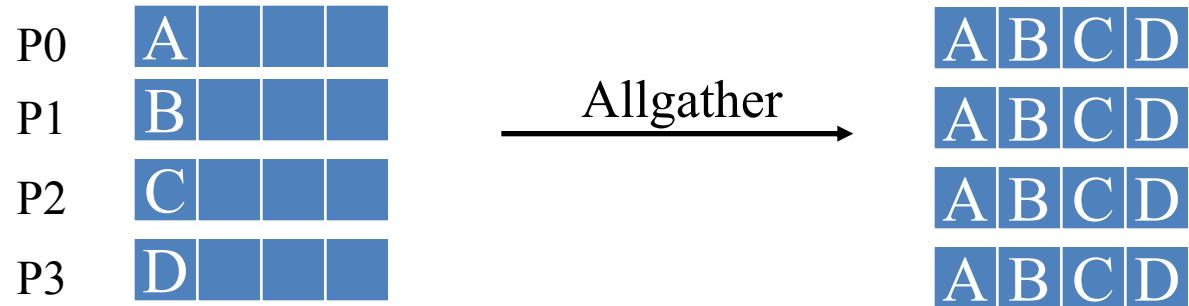
# Collective Operations in MPI

- Collective operations are called by all processes in a communicator
- **`MPI_BCAST`** distributes data from one process (the root) to all others in a communicator
- **`MPI_REDUCE`** combines data from all processes in communicator and returns it to one process

# Collective Data Movement



# More Collective Data Movement



# MPI References

- The Standard itself:
  - at <http://www mpi-forum.org>
  - All MPI official releases, in both postscript and HTML
- Other information on Web:
  - at <http://www.mcs.anl.gov/mpi>
  - pointers to lots of stuff, including other talks and tutorials, a FAQ, other MPI pages

# Programming Models

## Message Passing Models (MPI)

Fine-grained messages + computation

Hard to write code to deliver high coordination

Hard to maintain loadbalance

Hard to deal with disk locality, failures, stragglers

# “Cloud” Programming Models

- Data Parallel Models
- Restrict the programming interface
- Automatically handle failures, locality etc.

“Here’s an operation, run it on all of the data”

- I don’t care *where* it runs (you schedule that)
- In fact, feel free to run it *retry* on different nodes

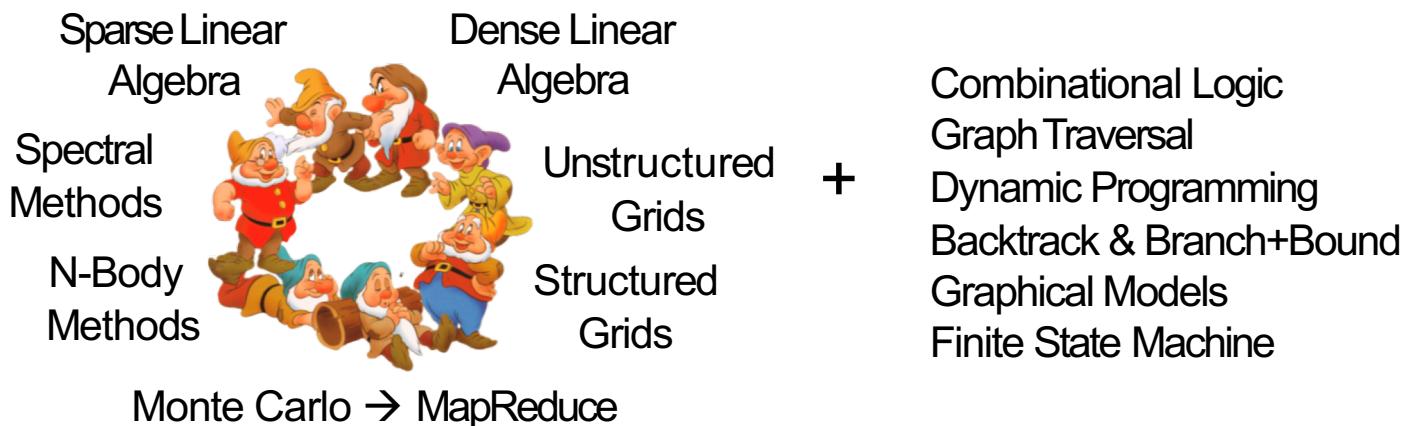
- Programming models: Map Reduce
- Frameworks and Engines: Hadoop , Spark, and many more.

# Seven/thirteen DwarFs/Patterns of Computing!



# What are the 13 Dwarfs?

- A (computational) dwarf is an algorithmic method that *captures a pattern of computation and communication*
  - Inspired by Phil Colella, who identified *seven* numerical methods important for science and engineering
- Benchmark Suite of 13 Computational Dwarfs & Apps

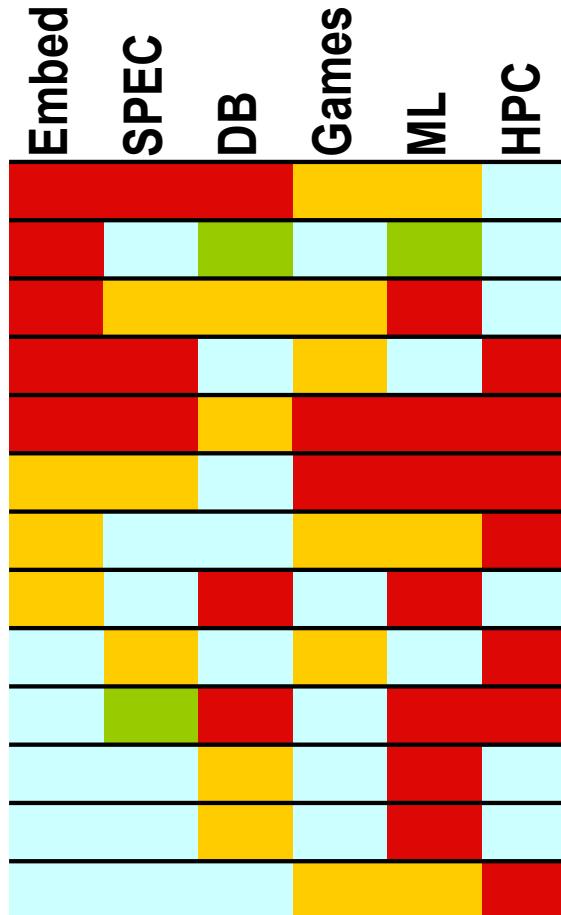


# What do commercial and CSE applications have in common?



Analyzed in detail in  
“Berkeley View” report

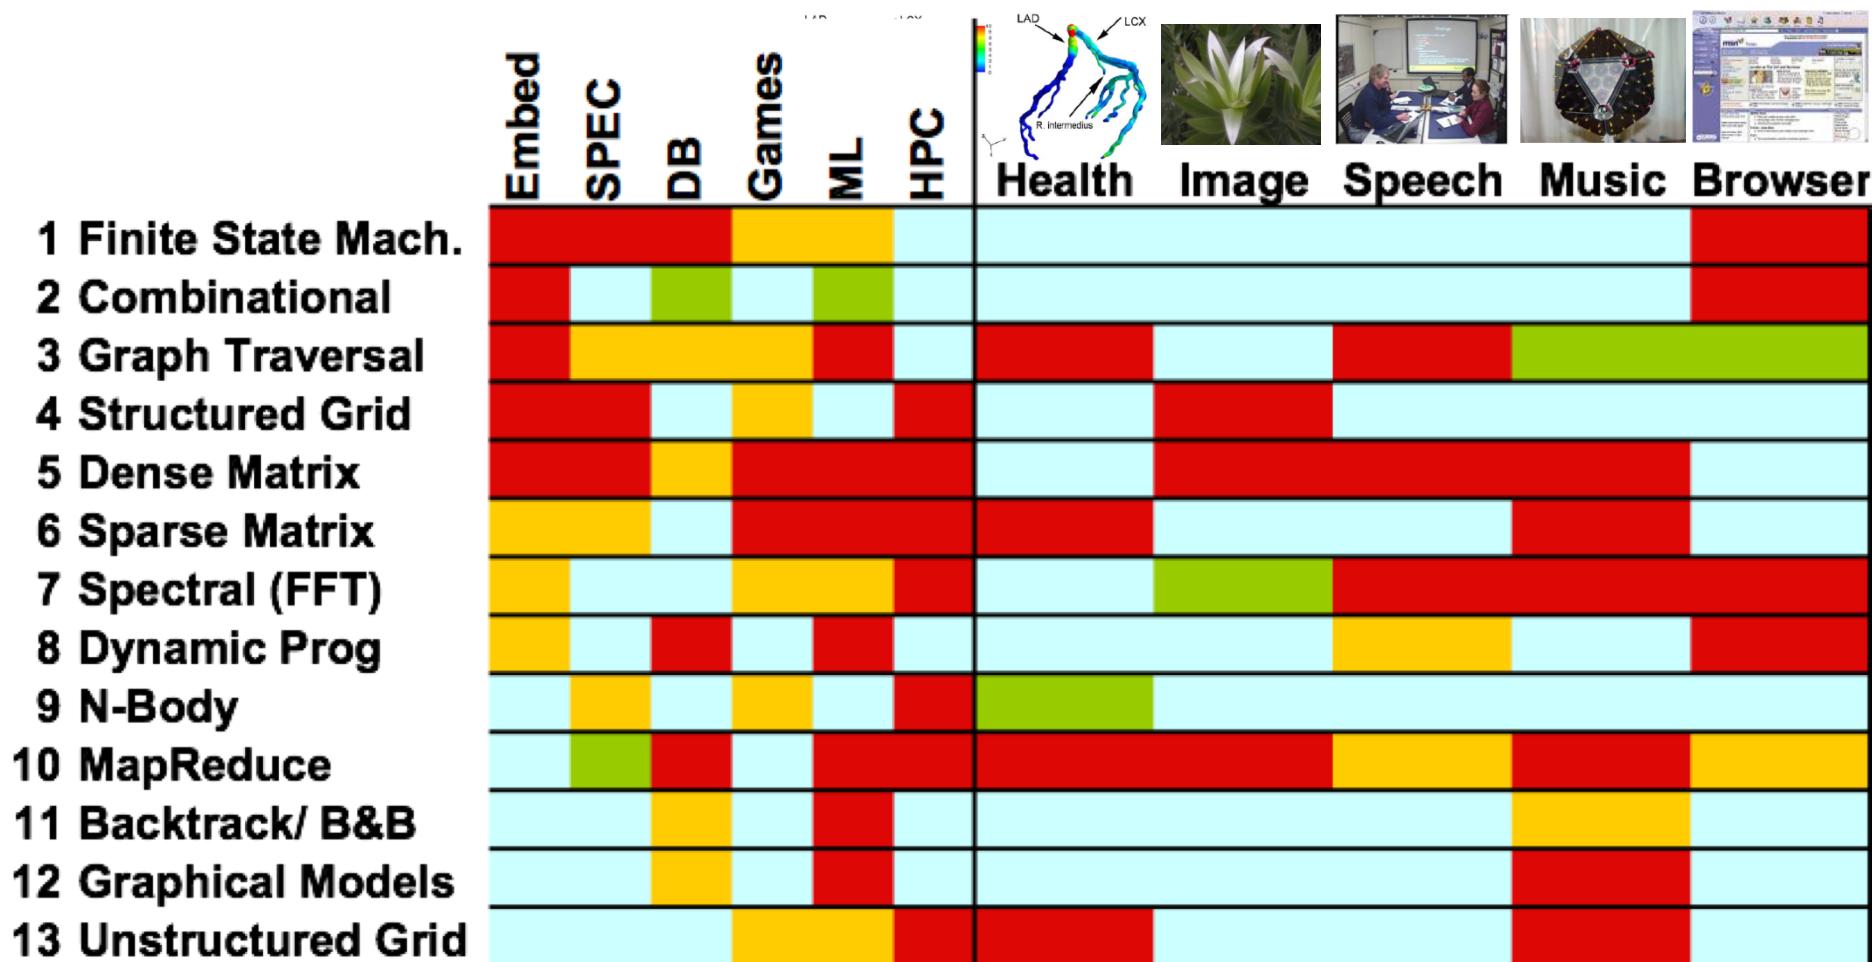
- 1 Finite State Mach.**
- 2 Combinational**
- 3 Graph Traversal**
- 4 Structured Grid**
- 5 Dense Matrix**
- 6 Sparse Matrix**
- 7 Spectral (FFT)**
- 8 Dynamic Prog**
- 9 N-Body**
- 10 MapReduce**
- 11 Backtrack/ B&B**
- 12 Graphical Models**
- 13 Unstructured Grid**



Analyzed in detail in  
“Berkeley View” report  
[www.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.html](http://www.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.html)

What do commercial and CSE applications have in common?

## Motif/Dwarf: Common Computational Methods (Red Hot → Blue Cool)



# N–Body Methods

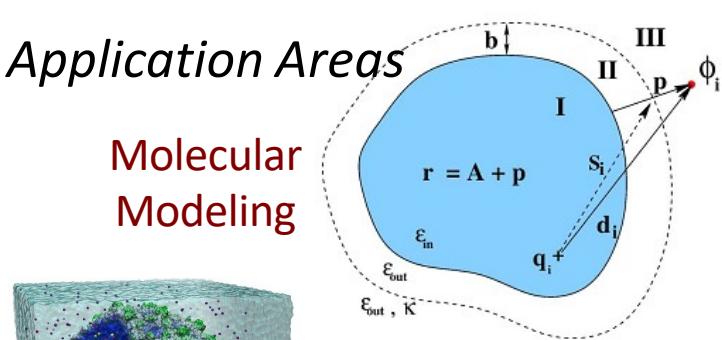
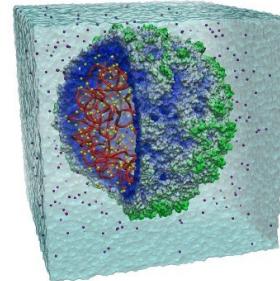
Calculations that depend on interactions between discrete points

## Approaches

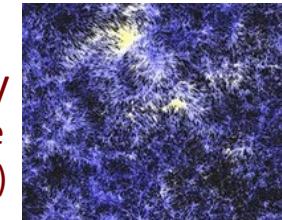
- In particle–particle methods
  - Every point depends on all others
    - $O(N^2)$ : Brute–force method
- Hierarchical particle methods
  - Combine forces or potentials from multiple points
    - $O(N \log N)$ : Barnes–Hutt
    - $O(N \log N)$ : Hierarchical charge partitioning
    - $O(N)$ : Fast multipole

## Application Areas

Molecular  
Modeling



Molecular  
Dynamics



Cosmology  
(Roadrunner Universe  
@ LANL)

# Sparse Linear Algebra

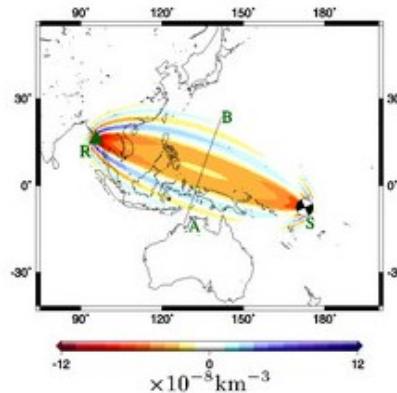
Multiplication involving matrices composed primarily of zeros

## Approaches

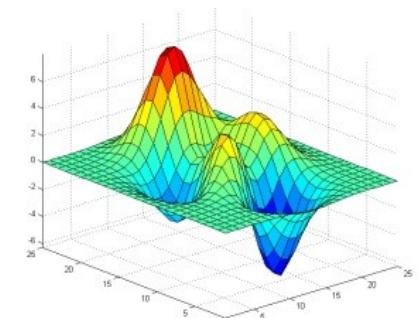
- Direct solving
  - Solve the problem with a series of operations
  - Large overhead
- Iterative methods
  - Use a series of successive approximations
    - Begins with a guess
    - Repeats until error is not significant

$$\begin{pmatrix} 0 & 3 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 7 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 7 \\ 0 & 0 & 0 & 5 & 0 & 0 & 0 & 8 & 0 & 0 \\ 2 & 0 & 0 & 0 & 0 & 0 & 6 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 9 & 0 & 0 & 0 & 4 & 0 \\ 0 & 1 & 0 & 7 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 4 \\ 0 & 0 & 7 & 0 & 0 & 0 & 0 & 5 & 0 & 0 \end{pmatrix}$$

## Application Areas & Kernels



Finite element analysis



Partial differential equations

# Dense Linear Algebra

Classic vector and matrix operations

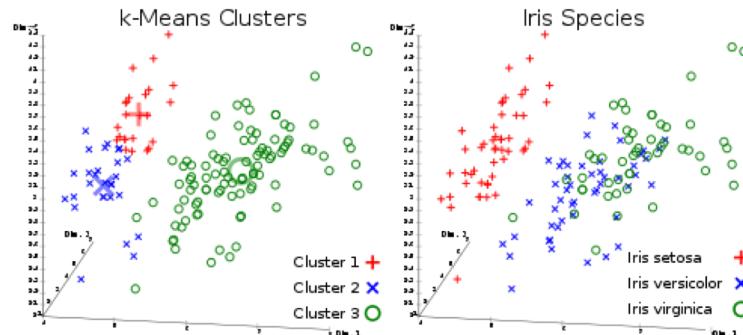
$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 2 & 3 & 4 & 5 \\ 6 & 7 & 8 & 9 \\ 0 & 1 & 2 & 3 \end{bmatrix} \times \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} = \begin{bmatrix} 30 \\ 40 \\ 80 \\ 20 \end{bmatrix}$$

## Approaches

- Uniprocessor Mapping
  - Naïve Method
    - Loops down rows & columns
  - Blocked Method
    - Makes use of spatial locality
- Parallel Mapping
  - 2-D Block Cyclic Distribution
    - Provides better data distribution and load balancing

## Application Areas

- Linear Algebra
  - LAPACK
  - ATLAS
- Data Mining
  - Streamcluster
  - K-means



Source: <http://view.eecs.berkeley.edu>

# Structured Grids

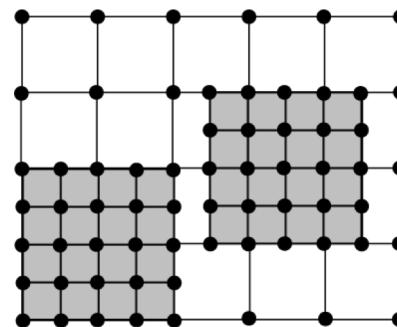
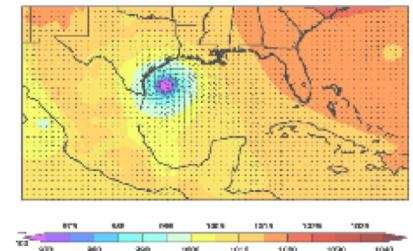
Computation steps update data in a regular multi-dimensional grid

## Approaches

- Regular Static Grid
  - Grid dimensions remain as they are from the start
    - High spatial locality, limited temporal locality
    - Parallel processors can be assigned subgrids
- Adaptive Mesh Refinement
  - Overlay higher resolution grids on areas of interest
    - Expensive boundary computations

## Application Areas

- Image Processing
  - SRAD
- Physics Simulations
  - HotSpot



# Unstructured Grids

Computations that depend on neighbors in an irregular grid

- *Approaches*

- In particle–particle methods
    - Every point depends on its neighbors
      - $O(N)$ : Brute–force method

- *Application Areas*

- Computational fluid dynamics
  - Belief propagation



# MapReduce

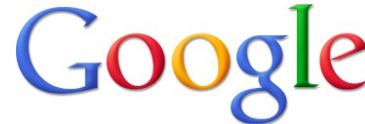
Process subsets of data independently and merge results



## Approaches

- Manual MapReduce
  - Application-specific “distribute” and “merge” functionality
    - mpiBLAST
- MapReduce Frameworks
  - Ease scaling of embarrassingly parallel applications
    - Hadoop and BOINC

## Application Areas

- Distributed Searching
  - Sequence Alignment
  - Parallel Monte Carlo Simulations
- 
- 

# Combinational Logic

Simple computation on large data sets, exhibiting bit-level parallelism

## Approaches

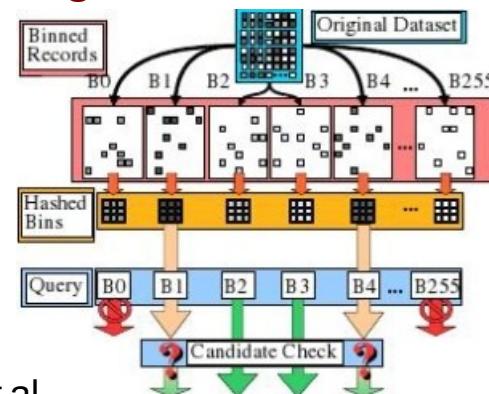
- Often naturally parallel
  - Can use pipelining techniques to increase throughput
  - $O(N)$ : Naive

## Application Areas

- Encryption & Decryption



- Hashing



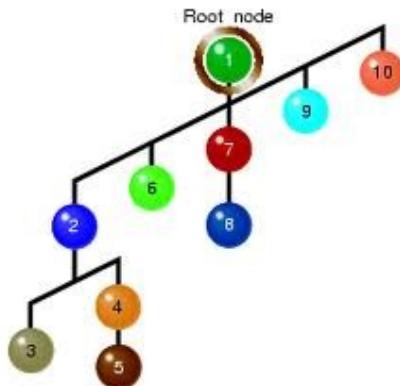
Source:  
John Owens et al.

# Graph Traversal

Traverse objects and examine them as they are traversed

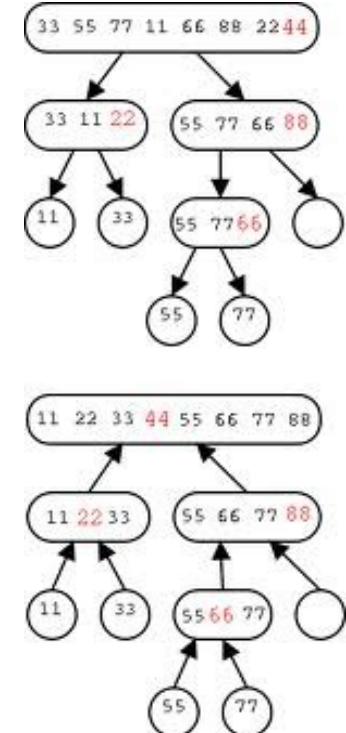
## Approaches

- Wavefront Movement
  - Traversal from single source
  - $O(N)$ : Brute-force method
- Divide and Conquer
  - Sorting Algorithms
  - $O(N \log N)$ : Quicksort



## Application Areas

- Searching
- Sorting
- Collision Detection



# Dynamic Programming

Compute solutions by solving simpler overlapping subproblems

## *Approaches*

- Uniprocessor Mapping
  - Solve subproblem once & store
    - Assumes fast store and lookup times of solutions
- Parallel Mapping
  - Solutions to overlapping subproblems communicated or re-computed
    - Trade-off between compute and communicate costs

## *Application Areas*

- Graph problems
  - Floyd's All-Pairs shortest path
  - Bellman-Ford algorithm
- Sequence alignment
  - Needleman-Wunsch
  - Smith-Waterman

# Backtracking and Branch & Bound

Branch-and-bound algorithms used to solve search and global optimization

## *Approaches*

- Heuristic
- Dynamic Load Balance
  - Divide search space (fairly) among available processors.

## *Application Areas*

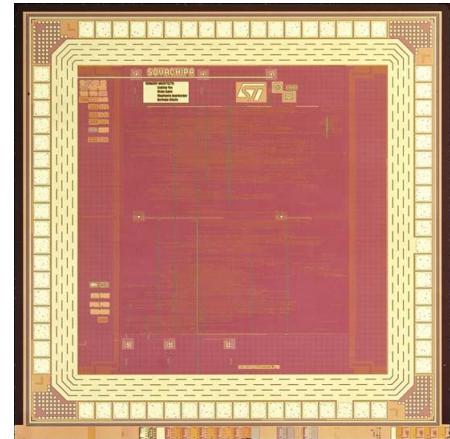
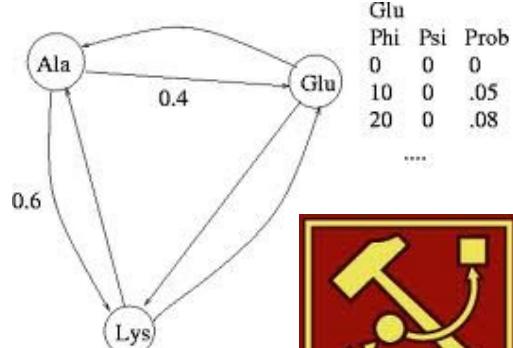
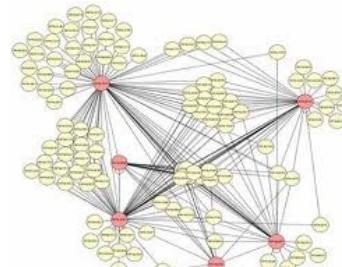
- Artificial Intelligence: N-Queens
- Integer Linear Programming
- Boolean Satisfiability
- Combinatorial Optimization

# Graphical Models

Construct graphs that represent random variables as nodes and conditional dependencies as edges

## *Application Areas*

- Computational Biology
  - Sequence homology search
- Machine Learning
  - Hidden Markov models
- Embedded Computing
  - Viterbi decode

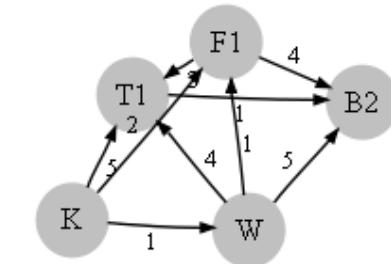
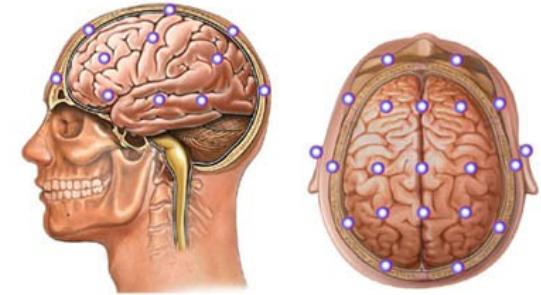
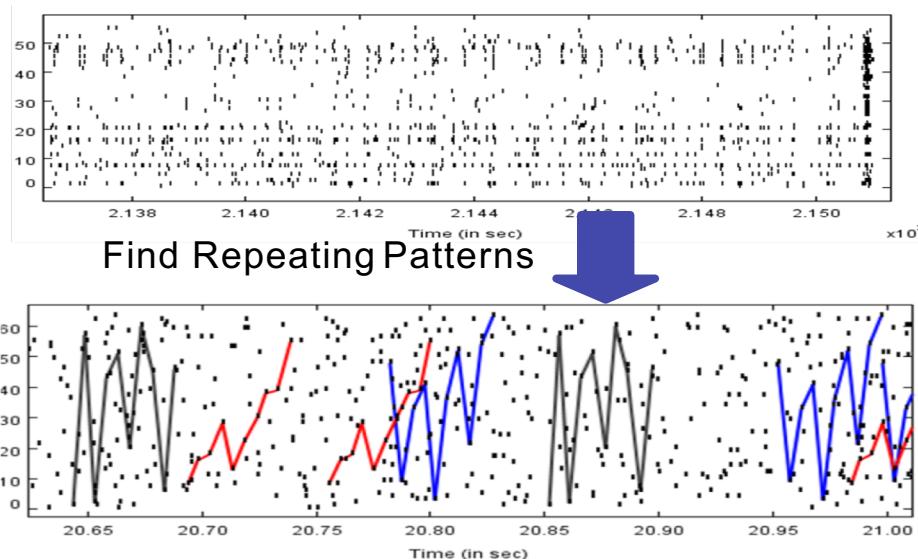


# Finite State Machine

Interconnected states which transition between one another

## Application Areas

- Video Decoding, Parsing, Compression
- Data Mining → Reverse Engineering the Brain



Infer  
Network Connectivity

"Temporal Data Mining for Neuroscience," *GPU Computing Gems*, Morgan Kaufmann, Feb. 2011.

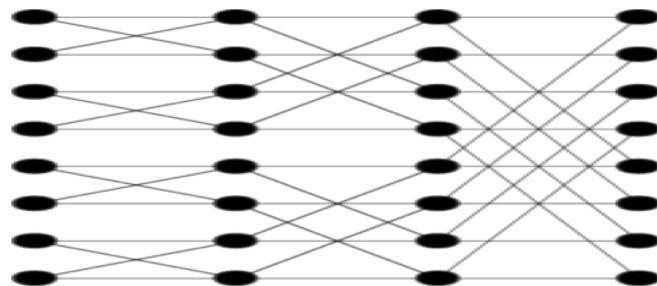
"High-Performance Biocomputing for Simulating the Spread of Contagion over Large Contact Networks," *BMC Genomics*, 2011.

# Spectral Methods

Spectral domain computations transformed from temporal or spatial domains and solved numerically

## Approach

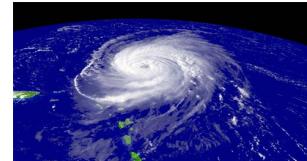
- ‘Butterfly’ pattern dependencies between stages of transformation
  - MAD operations on complex #
  - $O(N \log N)$ : FFT



Computational Organization

## Application Areas

### Fluid Dynamics



### Quantum Mechanics



### Weather Prediction

