

# STA2453 Lab 1

Yihan Duan

21/09/2021

## Exercise 1

### Data Quality

Load raw ED data.

```
knitr::opts_chunk$set(echo = TRUE)

# Load the required libraries
library(tidyverse)
library(lubridate)

# Read in the data
ed_data <- read_csv("raw_ed_data.csv")

# Show head
head(ed_data)
```

```
## # A tibble: 6 x 10
##   ENCOUNTER_NUM CTAS_CD CTAS_DESCR   ed_start_time   ed_end_time
##           <dbl> <chr>   <chr>         <dtm>          <dtm>
## 1             1 2      EMERGENCY 2019-01-01 06:06:00 2019-01-01 09:56:00
## 2             2 3      URGENT   2019-01-01 06:11:00 2019-01-01 06:43:00
## 3             3 3      URGENT   2019-01-01 06:21:00 2019-01-01 11:27:00
## 4             4 2      EMERGENCY 2019-01-01 06:36:00 2019-01-01 11:48:00
## 5             5 1      RESUSCITATION 2019-01-01 06:37:00 2019-01-01 08:27:00
## 6             6 2      EMERGENCY 2019-01-01 06:42:00 2019-01-01 11:48:00
## # ... with 5 more variables: ed_pia_time <dtm>, adm_start_time <dtm>,
## #   admitted <dbl>, los <dbl>, presenting_complaint <chr>
```

Clean up 'presenting\_complaint' column as shown in class.

```
# function to clean presenting complaints text
clean_complaints <- function(x) {

  x_clean <- x %>%
    # remove any leading and trailing spaces
    trimws() %>%
    # collapse > 1 blank space into 1 blank space
    gsub(" +", " ", .) %>%
    # set text to lower case
    tolower()

  return(x_clean)
```

```

}

ed_data <- ed_data %>%
  mutate(presenting_complaint = clean_complaints(presenting_complaint))

ed_data <- ed_data %>%
  mutate(presenting_complaint = case_when(presenting_complaint == "chest pian" ~
    "chest pain", presenting_complaint == "burns" ~ "burn", presenting_complaint ==
    "traumatic injuries" ~ "traumatic injury", presenting_complaint %in% c("unk",
    "missing") ~ "unknown", presenting_complaint == "headach" ~ "headache", TRUE ~
    presenting_complaint))

ed_data %>%
  count(presenting_complaint, sort = T) %>%
  mutate(proportion = round(n/sum(n), 3))

```

```

## # A tibble: 17 x 3
##   presenting_complaint      n proportion
##   <chr>                <int>      <dbl>
## 1 abdominal pain        14508      0.18
## 2 sore throat           13735      0.171
## 3 loss of hearing        8069      0.1
## 4 confusion              7742      0.096
## 5 headache               6979      0.087
## 6 upper extremity injury  4137      0.051
## 7 lower extremity injury  4077      0.051
## 8 back pain              3607      0.045
## 9 rash                   3490      0.043
## 10 chest pain            2767      0.034
## 11 general weakness      2764      0.034
## 12 traumatic injury      2610      0.032
## 13 hallucinations        2080      0.026
## 14 bizarre behaviour     1595      0.02
## 15 burn                  1232      0.015
## 16 trouble breathing      821      0.01
## 17 unknown               251      0.003

```

As shown in class, we can count all the 'NA's for each column.

```

count_NAs <- function(x) {

  num_NAs <- sum(is.na(x))

  return(num_NAs)
}

# the data
ed_data %>%
  # becomes the first argument passed to the summarize_all function
  summarize_all(count_NAs) %>%
  glimpse

```

```

## Rows: 1
## Columns: 10

```

```
## $ ENCOUNTER_NUM      <int> 0
## $ CTAS_CD             <int> 0
## $ CTAS_DESCR          <int> 0
## $ ed_start_time       <int> 793
## $ ed_end_time         <int> 396
## $ ed_pia_time         <int> 0
## $ adm_start_time      <int> 68849
## $ admitted           <int> 0
## $ los                 <int> 1186
## $ presenting_complaint <int> 0
```

```
ed_data %>%
  filter(is.na(ed_start_time) & is.na(ed_end_time))
```

```
## # A tibble: 3 x 10
##   ENCOUNTER_NUM CTAS_CD CTAS_DESCR   ed_start_time ed_end_time
##           <dbl> <chr>   <chr>         <dtm>         <dtm>
## 1         15461 4      SEMI-URGENT   NA            NA
## 2         29442 1      RESUSCITATION NA            NA
## 3         71909 3      URGENT       NA            NA
## # ... with 5 more variables: ed_pia_time <dtm>, adm_start_time <dtm>,
## #   admitted <dbl>, los <dbl>, presenting_complaint <chr>
```

Now we check all the variables one by one.

## ENCOUNTER\_NUM

First see if there are any duplicate numbers.

```
print("Length of ENCOUNTER_NUM:")
```

```
## [1] "Length of ENCOUNTER_NUM:"
```

```
length(ed_data$ENCOUNTER_NUM)
```

```
## [1] 80464
```

```
print("Number of unique values of ENCOUNTER_NUM:")
```

```
## [1] "Number of unique values of ENCOUNTER_NUM:"
```

```
length(unique(ed_data$ENCOUNTER_NUM))
```

```
## [1] 80248
```

We can see that there are duplicates in ENCOUNTER\_NUM column.

```
# Remove all duplicate rows
```

```
ed_data_dedup <- ed_data[!duplicated(ed_data), ]
```

Now check the number of unique values again.

```
print("Length of ENCOUNTER_NUM:")
```

```
## [1] "Length of ENCOUNTER_NUM:"
```

```
length(ed_data_dedup$ENCOUNTER_NUM)
```

```
## [1] 80249
```

```
print("Number of unique values of ENCOUNTER_NUM:")
```

```
## [1] "Number of unique values of ENCOUNTER_NUM:"
```

```
length(unique(ed_data_dedup$ENCOUNTER_NUM))
```

```
## [1] 80248
```

Still one duplicate ENCOUNTER\_NUM, lets find it.

```
num_freq <- ed_data_dedup %>%  
  count(ENCOUNTER_NUM) %>%  
  filter(n > 1)  
dup_num = num_freq$ENCOUNTER_NUM[1]  
ed_data_dedup %>%  
  filter(ENCOUNTER_NUM == dup_num)
```

```
## # A tibble: 2 x 10  
##   ENCOUNTER_NUM CTAS_CD CTAS_DESCR ed_start_time      ed_end_time  
##           <dbl> <chr>   <chr>      <dtm>          <dtm>  
## 1         44042 2      EMERGENCY 2019-07-18 16:38:00 2019-07-19 06:02:00  
## 2         44042 2      EMERGENCY 2019-07-18 16:38:00 2019-07-19 06:02:00  
## # ... with 5 more variables: ed_pia_time <dtm>, adm_start_time <dtm>,  
## #   admitted <dbl>, los <dbl>, presenting_complaint <chr>
```

Notice the only difference between these two records are the `ed_pia_time`.

## CTAS\_CD

Check data integrity for column 'CTAS\_CD'.

```
ed_data_dedup %>%  
  count(CTAS_CD)
```

```
## # A tibble: 6 x 2  
##   CTAS_CD      n  
##   <chr>   <int>  
## 1 1       3289  
## 2 2       26029  
## 3 3       34688  
## 4 4       12009  
## 5 5       3033  
## 6 N/A      1201
```

As we can see, the values presented in this column are mostly in range [1, 5]. However, there are many records with CTAS\_CD missing.

## CTAS\_DESCR

First check all the values presented in CTAS\_DESCR.

```
ed_data_dedup %>%  
  count(CTAS_DESCR)
```

```
## # A tibble: 6 x 2  
##   CTAS_DESCR      n  
##   <chr>      <int>  
## 1 EMERGENCY    26029  
## 2 N/A          1201  
## 3 NON URGENT   3033  
## 4 RESUSCITATION 3289
```

```
## 5 SEMI-URGENT 12009
## 6 URGENT      34688
```

There should be a 1 to 1 mapping from CTAS\_CD to CTAS\_DESCR.

```
ed_data[, c("CTAS_CD", "CTAS_DESCR")] %>%
  unique
```

```
## # A tibble: 6 x 2
##   CTAS_CD CTAS_DESCR
##   <chr>   <chr>
## 1 2      EMERGENCY
## 2 3      URGENT
## 3 1      RESUSCITATION
## 4 N/A    N/A
## 5 4      SEMI-URGENT
## 6 5      NON URGENT
```

Yes the mapping is 1 to 1.

ed\_start\_time, ed\_end\_time, ed\_pia\_time, adm\_start\_time

ed\_start\_time should always come before ed\_end\_time.

```
print("PIA before arrival at the ED:")
```

```
## [1] "PIA before arrival at the ED:"
```

```
ed_data %>%
  filter(ed_start_time > ed_pia_time) %>%
  nrow
```

```
## [1] 16
```

```
print("Departure before arrival at the ED:")
```

```
## [1] "Departure before arrival at the ED:"
```

```
ed_data %>%
  filter(ed_start_time > ed_end_time) %>%
  nrow
```

```
## [1] 1591
```

```
print("PIA after leaving the ED:")
```

```
## [1] "PIA after leaving the ED:"
```

```
ed_data %>%
  filter(ed_pia_time > ed_end_time) %>%
  nrow
```

```
## [1] 3934
```

```
print("PIA after admitted to the hospital:")
```

```
## [1] "PIA after admitted to the hospital:"
```

```
ed_data %>%
  filter(ed_pia_time > adm_start_time) %>%
  nrow
```

```
## [1] 441
```

```

print("Admitted to the hospital before leaving the ED:")

## [1] "Admitted to the hospital before leaving the ED:"
ed_data %>%
  filter(ed_end_time > adm_start_time) %>%
  nrow

## [1] 11297

print("Arrived at the ED after admitted to the hospital:")

## [1] "Arrived at the ED after admitted to the hospital:"
ed_data %>%
  filter(ed_start_time > adm_start_time) %>%
  nrow

## [1] 7

print("PIA time of '2099-01-01':")

## [1] "PIA time of '2099-01-01':"
ed_data_dedup %>%
  filter(ed_pia_time == ymd("2099-01-01")) %>%
  nrow

## [1] 1582

ed_data_dedup %>%
  filter(ed_pia_time == ymd("2099-01-01"))

## # A tibble: 1,582 x 10
##   ENCOUNTER_NUM CTAS_CD CTAS_DESCR ed_start_time      ed_end_time
##             <dbl> <chr>   <chr>      <dtm>          <dtm>
## 1             208 3      URGENT    2019-01-02 08:36:00 2019-01-02 21:32:00
## 2             228 3      URGENT    2019-01-02 13:26:00 2019-01-02 14:49:00
## 3             340 2      EMERGENCY 2019-01-02 21:00:00 2019-01-06 05:33:00
## 4             364 2      EMERGENCY 2019-01-02 22:33:00 2019-01-03 07:45:00
## 5             441 3      URGENT    2019-01-03 09:52:00 2019-01-03 19:24:00
## 6             478 2      EMERGENCY 2019-01-03 15:59:00 2019-01-03 17:52:00
## 7             529 3      URGENT    2019-01-03 19:29:00 2019-01-03 21:18:00
## 8             562 3      URGENT    2019-01-03 21:28:00 2019-01-04 01:48:00
## 9             598 2      EMERGENCY 2019-01-03 23:42:00 2019-01-04 10:00:00
## 10            641 2      EMERGENCY 2019-01-04 04:26:00 2019-01-03 14:35:00
## # ... with 1,572 more rows, and 5 more variables: ed_pia_time <dtm>,
## #   adm_start_time <dtm>, admitted <dbl>, los <dbl>,
## #   presenting_complaint <chr>

```

As we can see, many 'ed\_pia\_time' entries are labeled with unrealistic dates – '2099-01-01'.

### adm\_start\_time, admitted

adm\_start\_time only make sense if the patient is admitted.

```

ed_data %>%
  filter(admitted == 0) %>%
  count(adm_start_time)

```

```
## # A tibble: 1 x 2
##   adm_start_time      n
##   <dtm>          <int>
## 1 NA              68848
```

All patients that are not admitted have no 'adm\_start\_tme'.

```
ed_data %>%
  filter(admitted == 1) %>%
  count_NAs()
```

```
## [1] 371
```

There are 371 records that are admitted but have no 'adm\_start\_time'.

los

```
ed_data %>%
  filter(los == 24) %>%
  count
```


```
## # A tibble: 1 x 1
##       n
##   <int>
## 1  2772
```

## Descriptive analysis

### Length of stay

```
ed_data_dedup$startToEnd = as.numeric(difftime(ed_data_dedup$ed_end_time, ed_data_dedup$ed_start_time),
  units = "hours")

ed_data_dedup %>%
  filter(startToEnd >= 0) %>%
  ggplot(aes(x = startToEnd)) + geom_histogram(binwidth = 1) + labs(x = "(Departure Time - Arrival Time)",
    y = "Number of Records")
```



```
ggsave("los.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
ed_data_dedup %>%
```

```
  filter(startToEnd >= 0) %>%
```

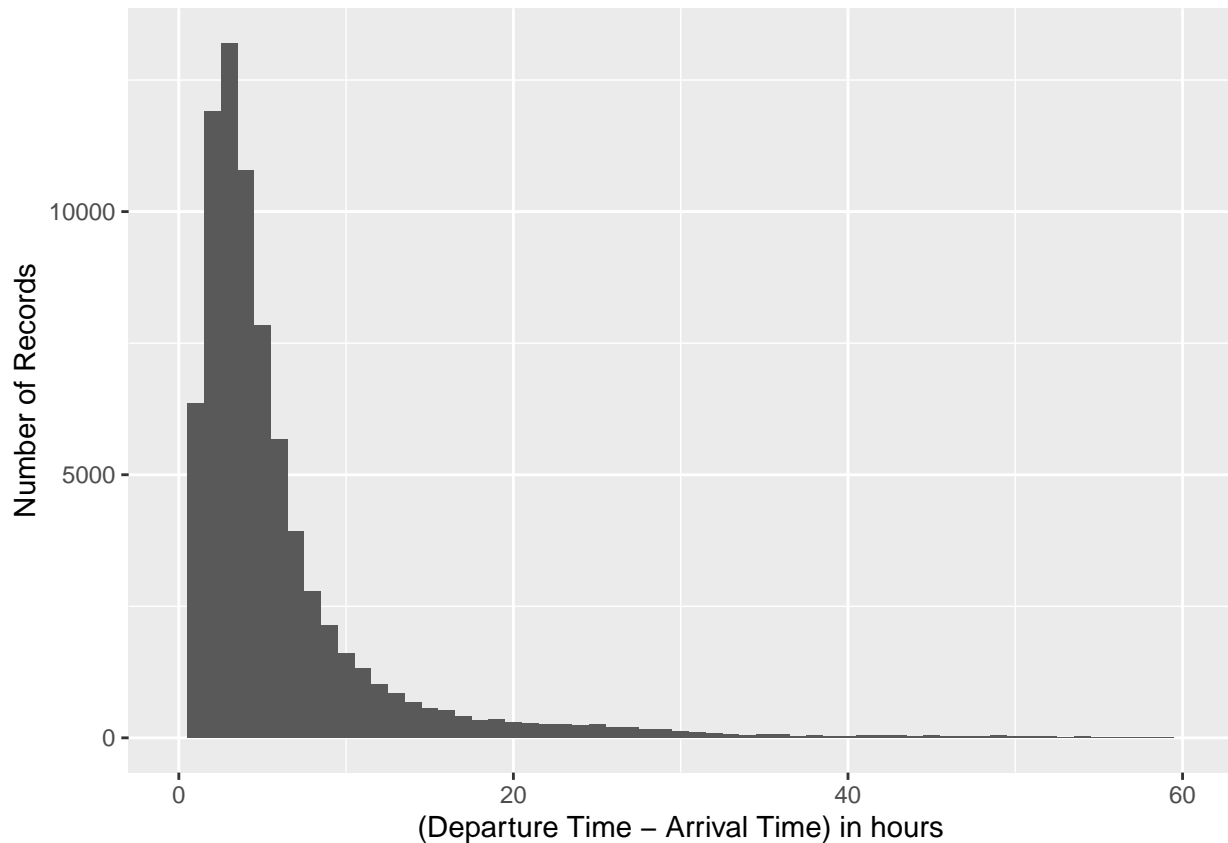
```
  ggplot(aes(x = startToEnd)) + geom_histogram(binwidth = 1) + xlim(0, 60) + labs(x = "(Departure Time - Arrival Time) in hours",
```

```
  y = "Number of Records")
```

```
## Warning: Removed 383 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```





```
ggsave("los_zoom.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
## Warning: Removed 383 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```

```
startToEndMean = mean(ed_data_dedup$startToEnd[ed_data_dedup$startToEnd > 0], na.rm = TRUE)
```

```
startToEndSd = sd(ed_data_dedup$startToEnd[ed_data_dedup$startToEnd > 0], na.rm = TRUE)
```

```
startToEndMean
```

```
## [1] 6.417152
```

```
startToEndSd
```

```
## [1] 9.136179
```

### Time to seeing a physician

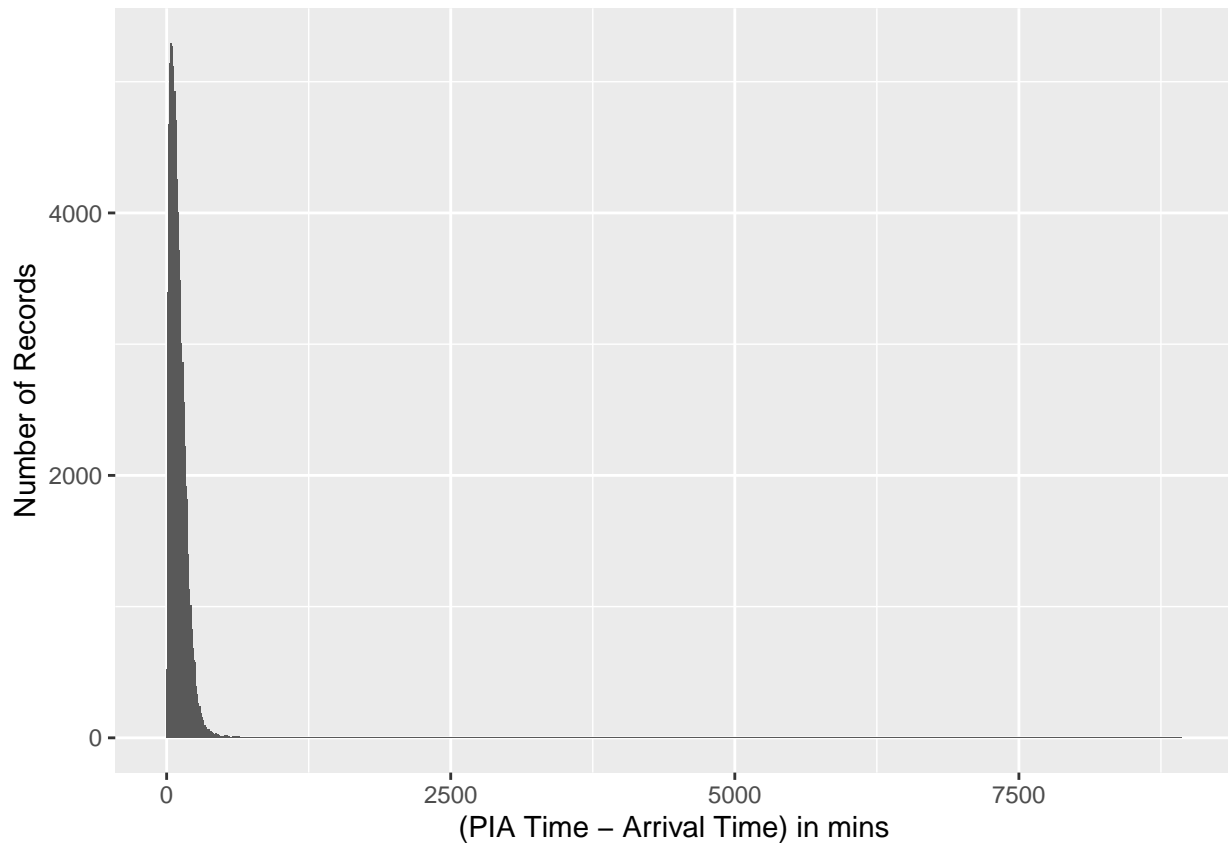
```
ed_data_dedup$startToPia = as.numeric(difftime(ed_data_dedup$ed_pia_time, ed_data_dedup$ed_start_time),
  units = "mins")
```

```
ed_data_dedup %>%
```

```
  filter(ed_pia_time != ymd("2099-01-01")) %>%
```

```
  filter(startToPia > 0) %>%
```

```
  ggplot(aes(x = startToPia)) + geom_histogram(binwidth = 10) + labs(x = "(PIA Time - Arrival Time) in minutes",
    y = "Number of Records")
```



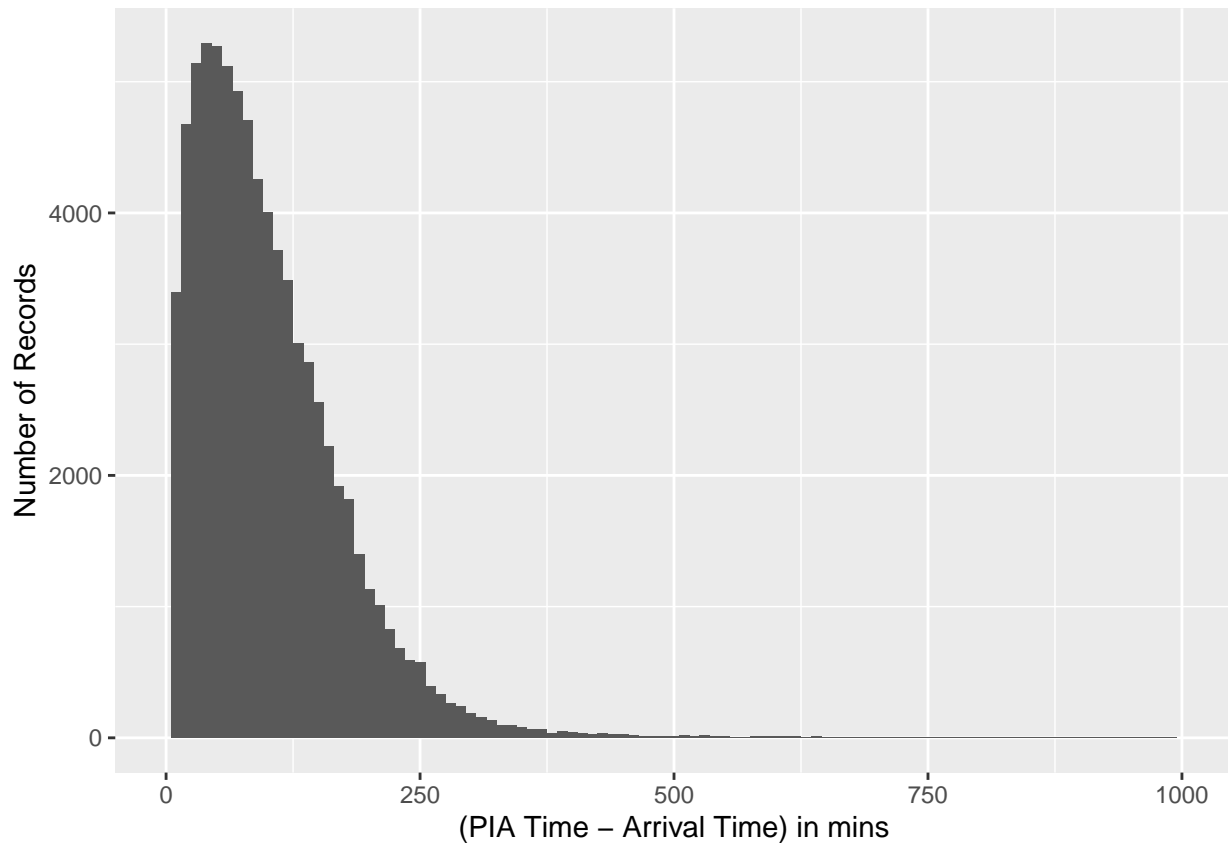
```
ggsave("pia.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
ed_data_dedup %>%  
  filter(ed_pia_time != ymd("2099-01-01")) %>%  
  filter(startToPia > 0) %>%  
  ggplot(aes(x = startToPia)) + geom_histogram(binwidth = 10) + xlim(0, 1000) +  
  labs(x = "(PIA Time - Arrival Time) in mins", y = "Number of Records")
```

```
## Warning: Removed 69 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```



```
ggsave("pia_zoom.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
## Warning: Removed 69 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```

```
startToPiaMean = mean(ed_data_dedup$startToPia[ed_data_dedup$startToPia > 0 & ed_data_dedup$ed_pia_time
  ymd("2099-01-01")], na.rm = TRUE)
startToPiaSd = sd(ed_data_dedup$startToPia[ed_data_dedup$startToPia > 0 & ed_data_dedup$ed_pia_time !=
  ymd("2099-01-01")], na.rm = TRUE)
```

```
startToPiaMean
```

```
## [1] 102.8916
```

```
startToPiaSd
```

```
## [1] 104.6482
```

### Counts and proportions of presenting complaints

```
PC_counts = ed_data_dedup %>%
  count(presenting_complaint) %>%
  arrange(desc(n))
```

```
jpeg(filename = "PC.jpg", width = 800, height = 600)
pie(PC_counts$n, labels = PC_counts$presenting_complaint, radius = 0.9, cex = 1.2)
```

```
dev.off()
```

```
## pdf  
## 2
```

```
PC_counts
```

```
## # A tibble: 17 x 2  
##   presenting_complaint      n  
##   <chr>                <int>  
## 1 abdominal pain        14470  
## 2 sore throat           13708  
## 3 loss of hearing        8052  
## 4 confusion              7718  
## 5 headache               6948  
## 6 upper extremity injury  4131  
## 7 lower extremity injury  4064  
## 8 back pain              3600  
## 9 rash                   3478  
## 10 general weakness      2760  
## 11 chest pain            2755  
## 12 traumatic injury      2606  
## 13 hallucinations        2075  
## 14 bizarre behaviour     1587  
## 15 burn                  1229  
## 16 trouble breathing      817  
## 17 unknown               251
```

## CTAS

```
CTAS_counts = ed_data_dedup %>%  
  count(CTAS_DESCR) %>%  
  arrange(desc(n))  
  
jpeg(filename = "CTAS.jpg", width = 800, height = 600)  
pie(CTAS_counts$n, labels = CTAS_counts$CTAS_DESCR, radius = 1, cex = 1)  
dev.off()
```

```
## pdf  
## 2
```

```
CTAS_counts
```


```
## # A tibble: 6 x 2  
##   CTAS_DESCR      n  
##   <chr>        <int>  
## 1 URGENT       34688  
## 2 EMERGENCY    26029  
## 3 SEMI-URGENT  12009  
## 4 RESUSCITATION 3289  
## 5 NON URGENT   3033  
## 6 N/A         1201
```

## Number of encounters

```
ed_data_dedup$ed_start_time_YMD = as_date(ed_data_dedup$ed_start_time)

n_encounters = count(ed_data_dedup, ed_start_time_YMD)

ed_data_dedup %>%
  count(ed_start_time_YMD) %>%
  ggplot(aes(x = n)) + geom_histogram(binwidth = 5) + labs(x = "Number of Encounters",
  y = "Count")
```



```
ggsave("n_enc.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
n_encountersMean = mean(n_encounters$n, rm.na=TRUE)
```

```
n_encountersSd = sd(n_encounters$n)
```

```
n_encounters %>%
```

```
  ggplot(aes(x=n)) +
```

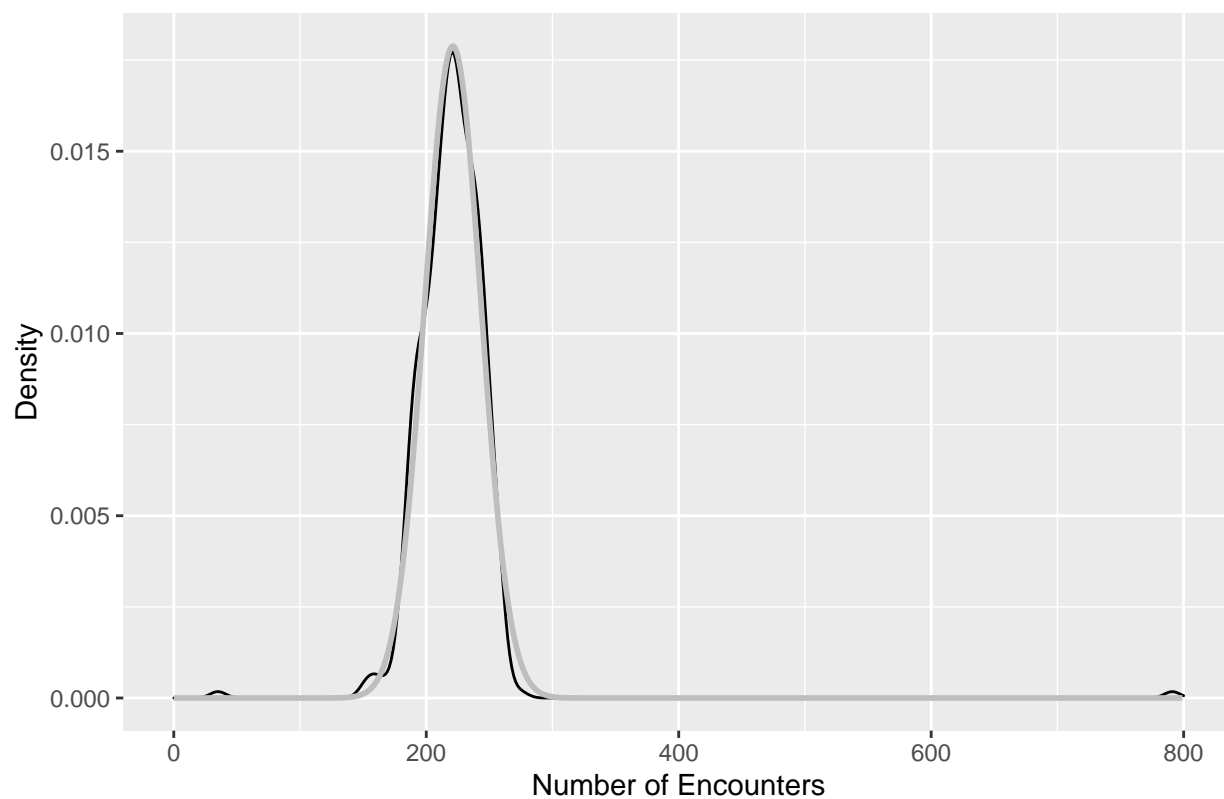
```
  geom_density() +
```

```
  stat_function(fun=dnorm, args=c(n_encountersMean, n_encountersSd / 1.7), xlim=c(0, 799), n=800, size=
```

```
  xlim(0,800) +
```

```
  labs(x = "Number of Encounters", y="Density", title="Real Density vs. Normal(221, 22)")
```

### Real Density vs. Normal(221, 22)



```
ggsave("n_enc_fit.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
n_encountersMean
```

```
## [1] 221.0716
```


```
n_encountersSd
```

```
## [1] 37.91762
```

### Time series

```
ggplot(n_encounters, aes(x = ed_start_time_YMD, y = n)) + geom_line() + labs(x = "Date",  
  y = "ED arrivals per day", title = "ED arrivals per day")
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```



```
ggsave("ed_arr.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
n_encountersMean
```

```
## [1] 221.0716
```

```
tail(n_encounters)
```

```
## # A tibble: 6 x 2
##   ed_start_time_YMD      n
##   <date>             <int>
## 1 2019-12-28         228
## 2 2019-12-29         224
## 3 2019-12-30         208
## 4 2019-12-31         231
## 5 2020-01-01          35
## 6 NA                791
```

### Weekend vs. weekday volumes

```
weekend_names = c("Sat", "Sun")
```

```
# Get weekday and weekend dfs
```

```
weekday_encounters = n_encounters %>%
  filter(!weekdays(ed_start_time_YMD, abbreviate = TRUE) %in% weekend_names) %>%
```

```

    filter(!is.na(ed_start_time_YMD))
weekend_encounters = n_encounters %>%
    filter(weekdays(ed_start_time_YMD, abbreviate = TRUE) %in% weekend_names)

# Get weekday stats
print("Weekday mean:")

## [1] "Weekday mean:"
mean(weekday_encounters$n)

## [1] 222.0426
print("Weekday standard deviation:")

## [1] "Weekday standard deviation:"
sd(weekday_encounters$n)

## [1] 24.28034
# Get weekend stats
print("Weekend mean:")

## [1] "Weekend mean:"
mean(weekend_encounters$n)

## [1] 213.1827
print("Weekend standard deviation:")


## [1] "Weekend standard deviation:"
sd(weekend_encounters$n)

## [1] 19.06034
weekday_encounters$day = "Weekdays"
weekend_encounters$day = "Weekends"
all_encounters = rbind(weekday_encounters, weekend_encounters)

all_encounters %>%
  ggplot(aes(x = day, y = n)) + geom_boxplot() + labs(x = "Weekday vs. weekend volumes",
    y = "ED arrivals perday")

```





```
ggsave("wkd_wke.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
t.test(n ~ day, data = all_encounters, var.equal = TRUE)
```

```
##
## Two Sample t-test
##
## data: n by day
## t = 3.3297, df = 360, p-value = 0.0009593
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  3.627101 14.092786
## sample estimates:
## mean in group Weekdays mean in group Weekends
##           222.0426           213.1827
```

## Working hours

```
ed_data_clean <- ed_data_dedup %>%
  filter(!is.na(ed_start_time), !is.na(ed_end_time)) %>%
  filter(ed_start_time < ed_end_time) %>%
  mutate(ed_start_time = floor_date(ed_start_time, unit = "hour"), ed_end_time = floor_date(ed_end_time,
    unit = "hour"))

arrivals <- ed_data_clean %>%
```

```

    select(timestamp = ed_start_time) %>%
    mutate(counter = 1)

departures <- ed_data_clean %>%
  select(timestamp = ed_end_time) %>%
  mutate(counter = -1)

census_volumes <- arrivals %>%
  bind_rows(departures) %>%
  arrange(timestamp, counter) %>%
  mutate(volume = cumsum(counter))

start <- min(ed_data_clean$ed_start_time)
end <- max(ed_data_clean$ed_end_time)
full_time_window <- tibble(timestamp = seq(start, end, by = "hours"))


census_volumes <- census_volumes %>%
  right_join(full_time_window, by = "timestamp") %>%
  arrange(timestamp) %>%
  fill(volume, .direction = "down")

census_volumes <- census_volumes %>%
  arrange(timestamp, volume) %>%
  group_by(timestamp) %>%
  summarise_all(last)

census_volumes %>%
  ggplot(aes(timestamp, volume)) + geom_line() + labs(x = "Date", title = "Emergency Department Census")

```

## Emergency Department Census Throughout 2019




```
ggsave("census_ydm_hms.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
census_hours <- census_volumes %>%  
  mutate(timestamp = format(timestamp, format = "%T"))  
  
aggregate(x = census_hours$volume, by = list(timestamp = census_hours$timestamp),  
  FUN = mean) %>%  
  ggplot(aes(timestamp, x, group = 1)) + geom_line() + labs(y = "Volume", x = "Time of Day",  
    title = "Average ED Census by Hour") + theme(axis.text.x = element_text(angle = 45))
```

### Average ED Census by Hour



```
ggsave("census_hms.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
census_by_hour = aggregate(x = census_hours$volume, by =list(timestamp=census_hours$timestamp), FUN=mean)
```

```
summary(census_by_hour)
```

```
##   timestamp      x
## Length:24      Min.   :40.14
## Class :character 1st Qu.:45.88
## Mode  :character Median :60.47
##                  Mean   :57.99
##                  3rd Qu.:69.25
##                  Max.   :73.24
```

### Descriptive Summary

```
ed_data_dedup %>%
  filter(startToEnd >= 0) %>%
  filter(ed_pia_time != ymd("2099-01-01")) %>%
  filter(startToPia >= 0) %>%
  summary()
```

```
## ENCOUNTER_NUM    CTAS_CD      CTAS_DESCR
## Min.   :      1  Length:75940  Length:75940
## 1st Qu.:20930    Class :character  Class :character
## Median :41048    Mode  :character  Mode  :character
```

```

## Mean      :40826
## 3rd Qu.   :61159
## Max.      :81133
##
## ed_start_time      ed_end_time
## Min.      :2019-01-01 06:06:00 Min.      :2019-01-01 06:43:00
## 1st Qu.   :2019-04-06 05:36:30 1st Qu.   :2019-04-06 16:34:45
## Median    :2019-07-05 04:24:30 Median    :2019-07-05 13:19:00
## Mean      :2019-07-03 21:40:02 Mean      :2019-07-04 04:04:32
## 3rd Qu.   :2019-10-01 19:53:00 3rd Qu.   :2019-10-02 01:55:45
## Max.      :2020-01-01 05:38:00 Max.      :2020-01-04 03:08:00
##
## ed_pia_time      adm_start_time      admitted
## Min.      :2019-01-01 06:22:00 Min.      :2019-01-01 07:56:00 Min.      :0.000
## 1st Qu.   :2019-04-06 07:41:45 1st Qu.   :2019-03-31 22:23:15 1st Qu.   :0.000
## Median    :2019-07-05 07:47:00 Median    :2019-07-06 20:37:30 Median    :0.000
## Mean      :2019-07-03 23:23:02 Mean      :2019-07-04 02:20:40 Mean      :0.144
## 3rd Qu.   :2019-10-01 21:54:00 3rd Qu.   :2019-10-03 06:26:15 3rd Qu.   :0.000
## Max.      :2020-01-01 09:21:00 Max.      :2020-01-01 11:48:00 Max.      :1.000
##
## NA's      :65006
##      los      presenting_complaint      startToEnd      startToPia
## Min.      : 0.100      Length:75940      Min.      : 0.100      Min.      : 3
## 1st Qu.   : 2.500      Class :character      1st Qu.   : 2.500      1st Qu.   : 46
## Median    : 4.017      Mode  :character      Median    : 4.017      Median    : 85
## Mean      : 5.780                                Mean      : 6.408      Mean      : 103
## 3rd Qu.   : 6.750                                3rd Qu.   : 6.750      3rd Qu.   : 139
## Max.      :24.000                                Max.      :169.550      Max.      :8935
##
## ed_start_time_YMD
## Min.      :2019-01-01
## 1st Qu.   :2019-04-06
## Median    :2019-07-05
## Mean      :2019-07-03
## 3rd Qu.   :2019-10-01
## Max.      :2020-01-01
##

```